

Sarthi AI Agent Development Platform (SADP)

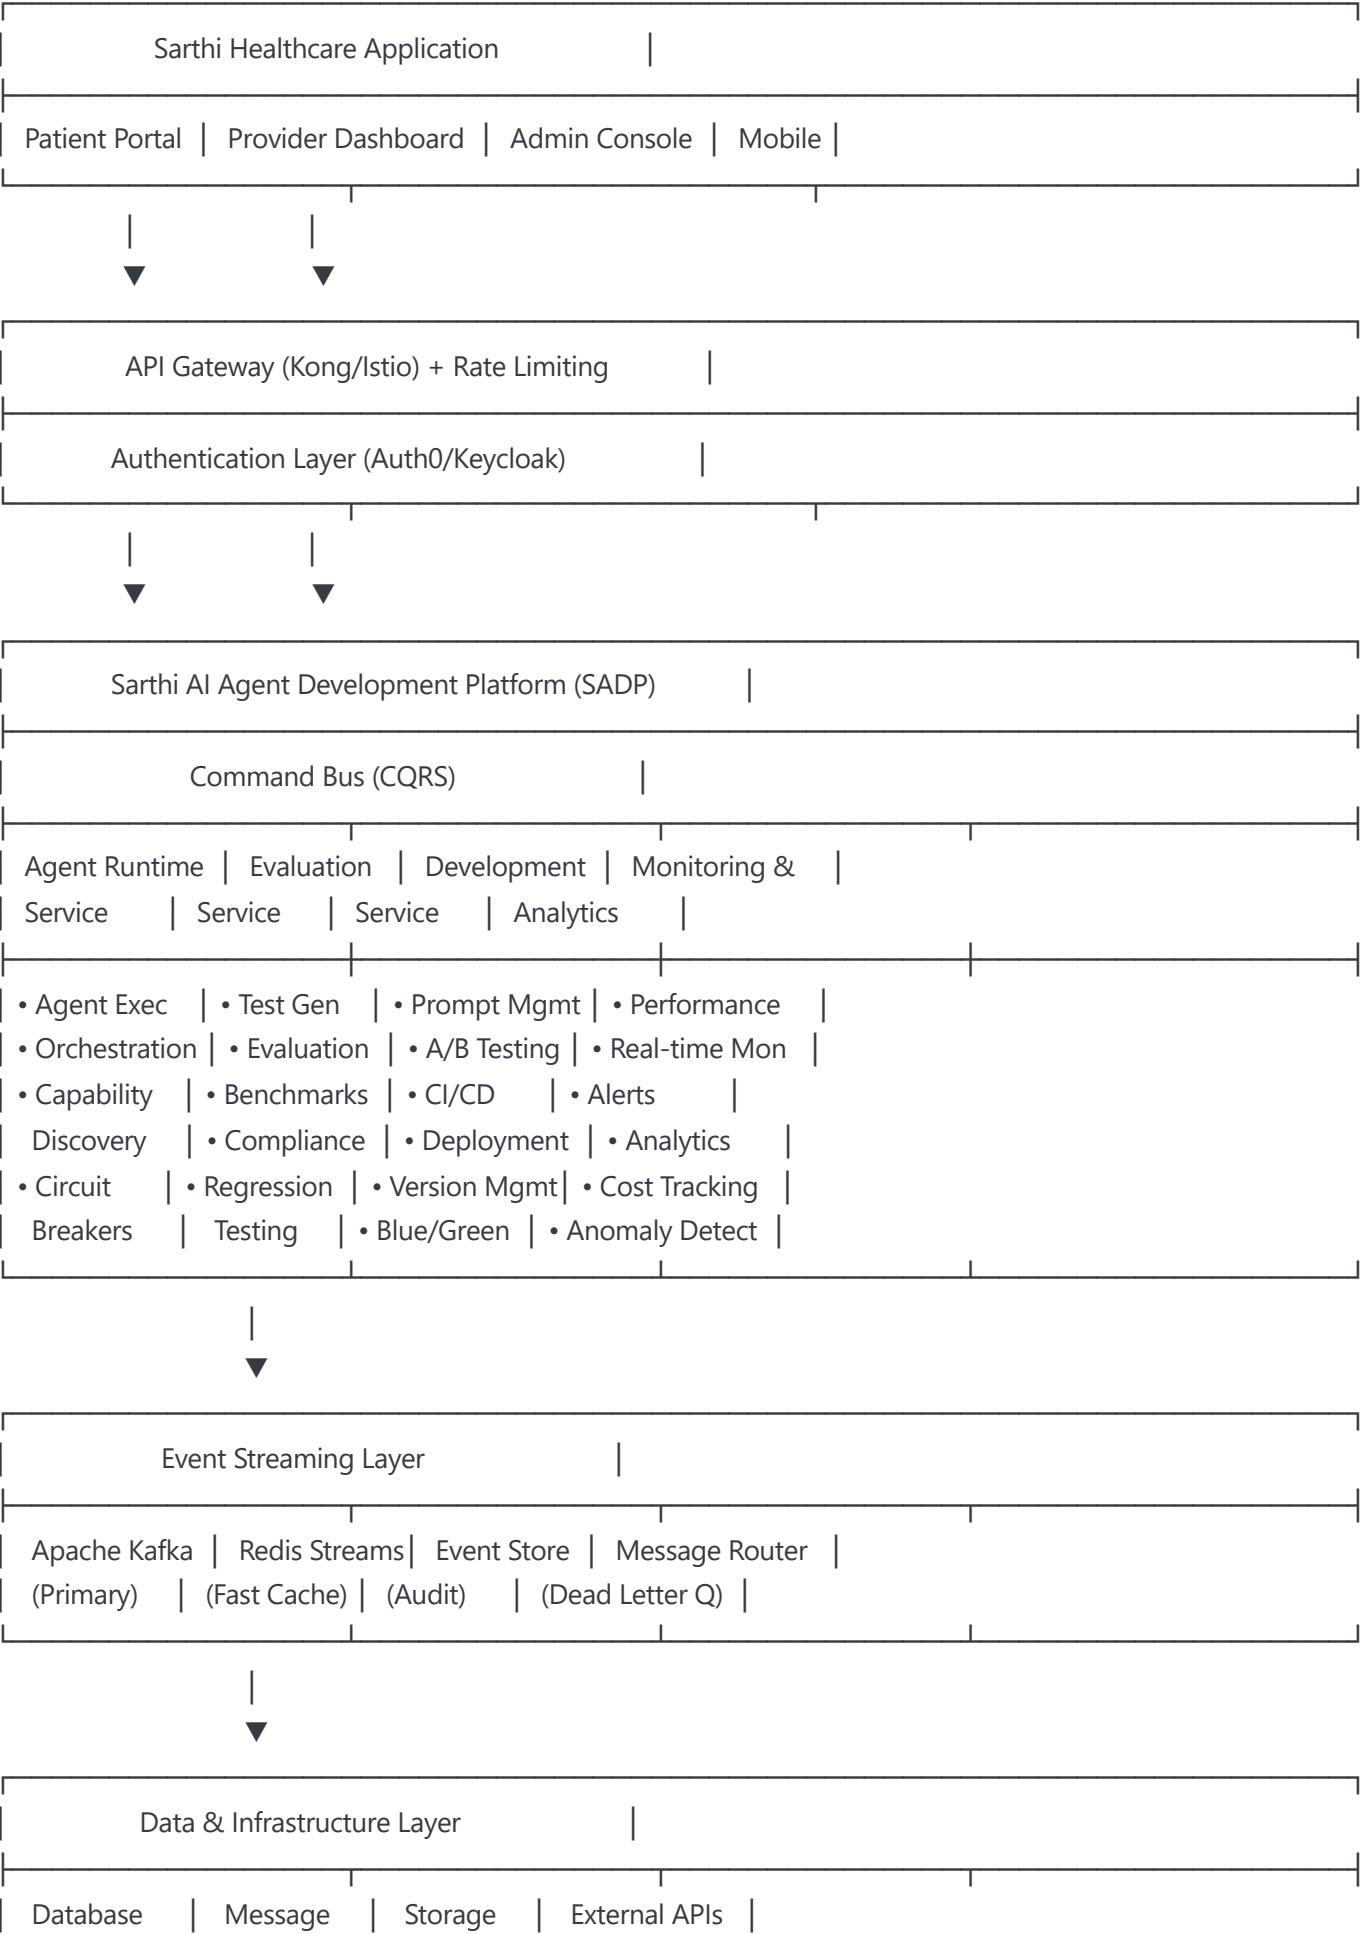
Production-Ready Independent Service Architecture

Overview

The Sarthi AI Agent Development Platform (SADP) is a standalone, enterprise-grade microservice that provides comprehensive AI agent development, evaluation, and management capabilities as an external API service for the main Sarthi healthcare application.

Architectural Improvements & Enhancements

1. Event-Driven Architecture with CQRS



Cluster	Queue	Services	& Integrations	
• PostgreSQL (Primary)	• Redis Cluster	• GCS Buckets • File System	• Claude API • OpenAI API	
• CockroachDB (Distributed)	• Kafka • RabbitMQ	• CDN • Object Store	• Healthcare APIs • FHIR Services	
• MongoDB (Documents)	(Fallback) • NATS	• Backup Storage	• EHR Systems • Monitoring SaaS	
• InfluxDB (Time Series)	(Streaming)	• Artifact Registry	• Observability	Stack

2. Enhanced Service Mesh Architecture

Service Discovery & Communication

```
yaml
```

Istio Service Mesh Configuration

apiVersion: networking.istio.io/v1beta1

kind: VirtualService

metadata:

name: sadp-routing

spec:

hosts:

- sadp-api

http:

- match:

- headers:

- version:

- exact: "v2"

route:

- destination:

- host: sadp-api

- subset: v2

- weight: 20

- destination:

- host: sadp-api

- subset: v1

- weight: 80

- route:

- destination:

- host: sadp-api

- subset: v1

Circuit Breaker Configuration

apiVersion: networking.istio.io/v1beta1

kind: DestinationRule

metadata:

name: sadp-circuit-breaker

spec:

host: sadp-api

trafficPolicy:

circuitBreaker:

- consecutiveErrors: 5

- interval: 30s

- baseEjectionTime: 30s

- maxEjectionPercent: 50

retryPolicy:

attempts: 3

perTryTimeout: 2s

Sarthi AI Agent Development Platform (SADP)

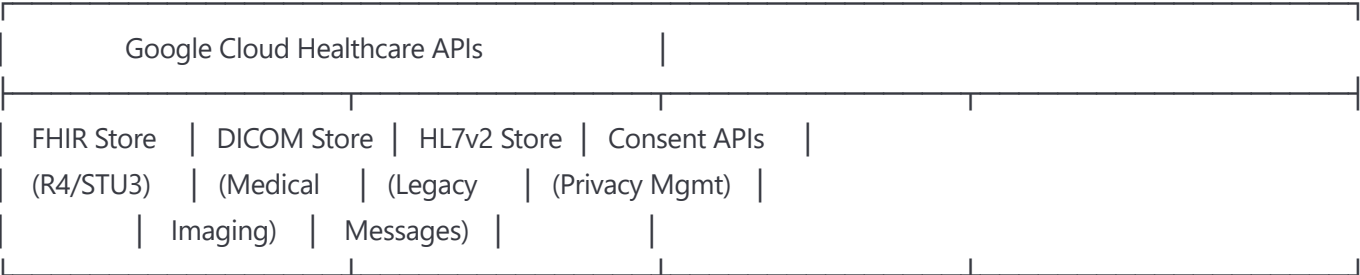
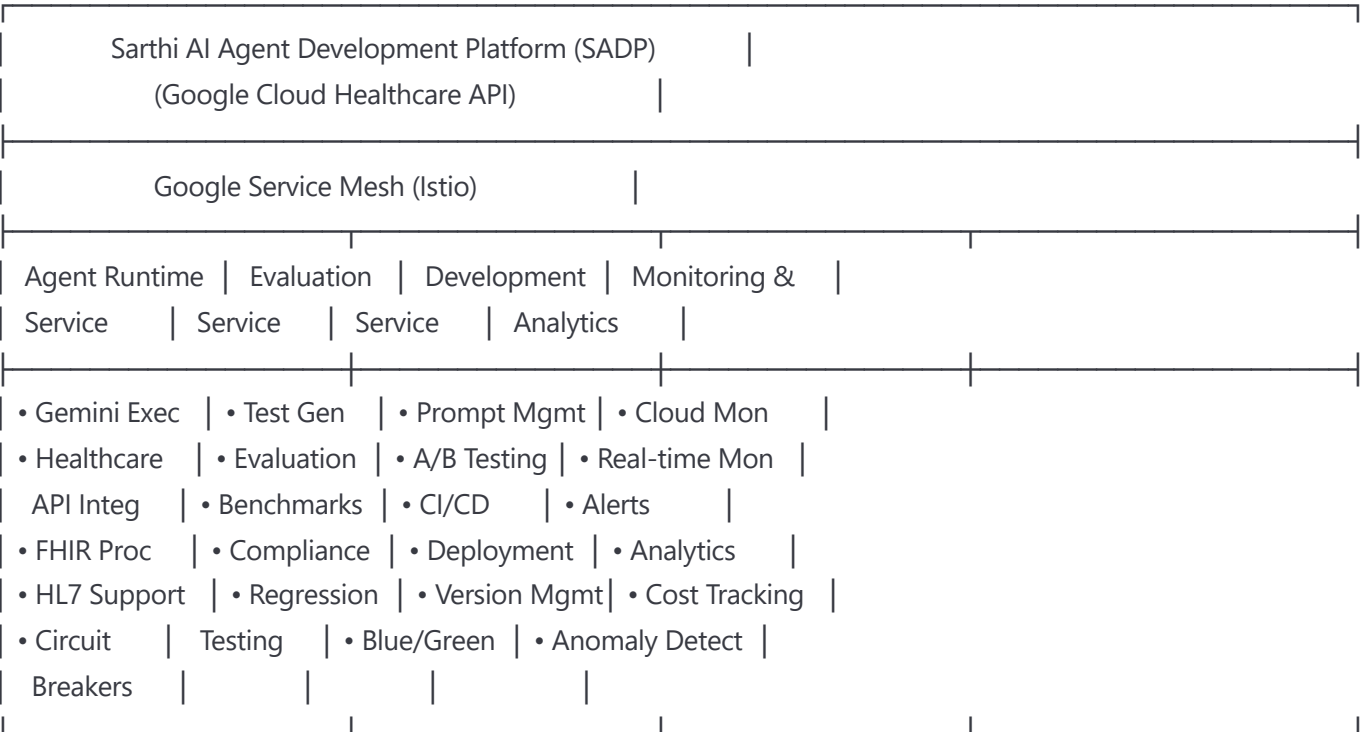


Google Healthcare Platform Integration Architecture

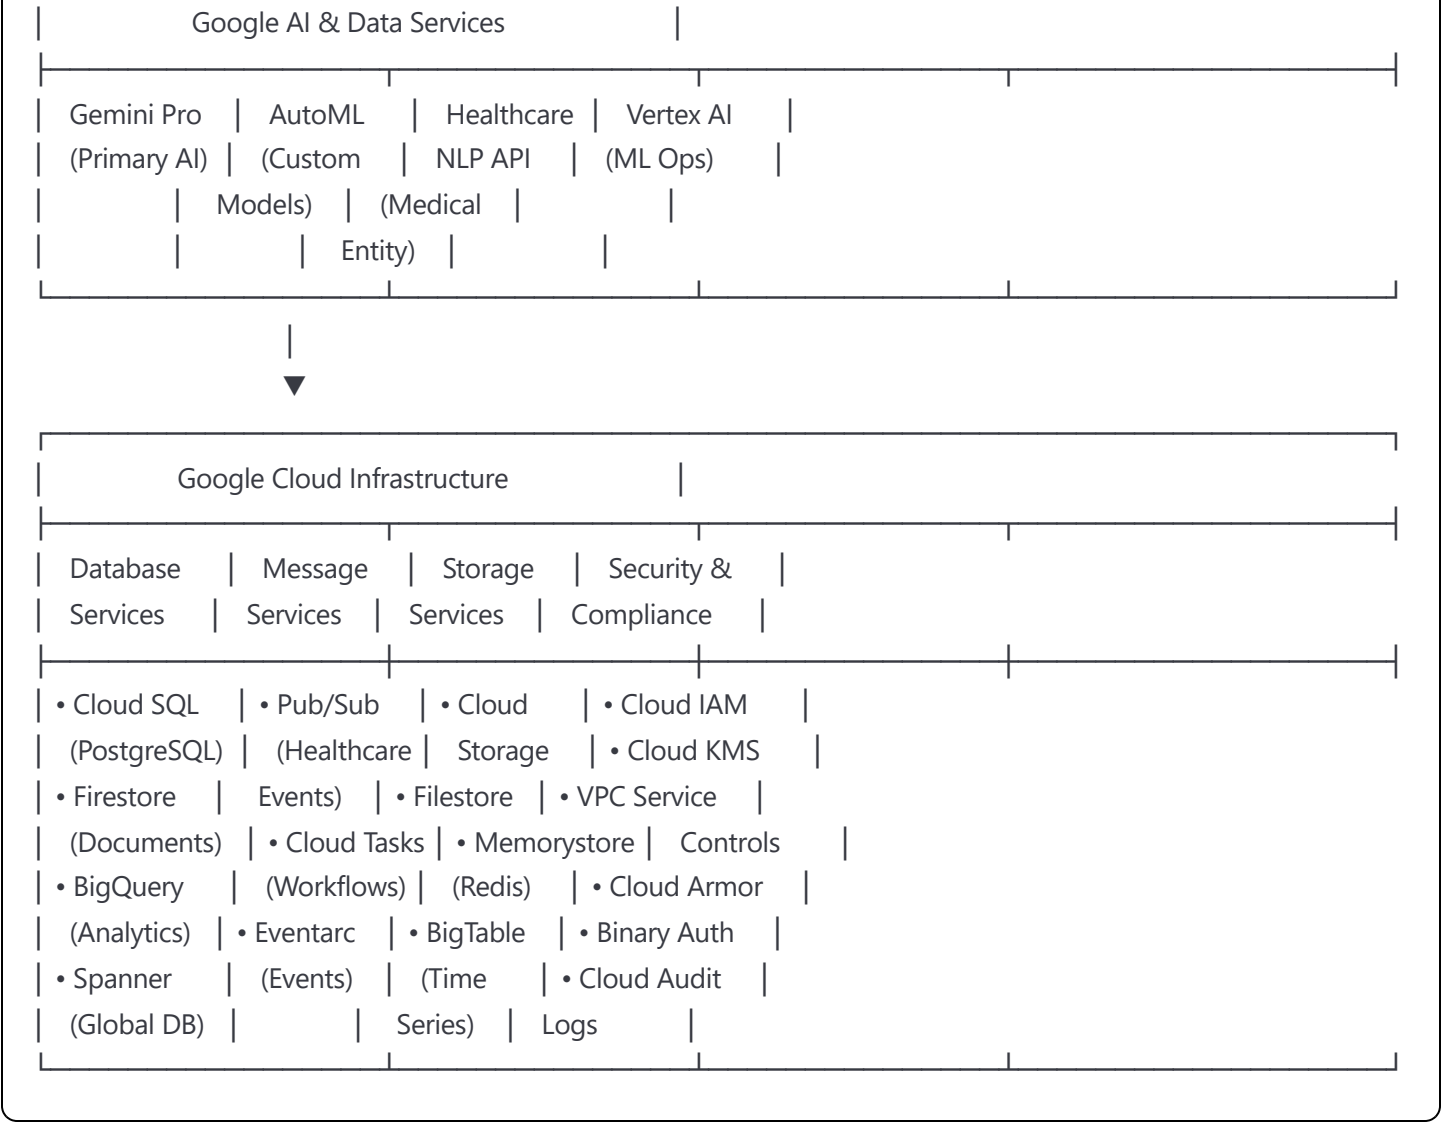
Overview

The Sarthi AI Agent Development Platform (SADP) is a standalone, enterprise-grade microservice built specifically for **Google Cloud Healthcare APIs** and **Google AI (Gemini)** models. It provides comprehensive AI agent development, evaluation, and management capabilities as an external API service for the main Sarthi healthcare application, leveraging Google's healthcare-specific infrastructure and AI capabilities.

Google Cloud Healthcare Architecture

High-Level Architecture





Google Cloud Healthcare Platform Integration

1. Google Cloud Healthcare APIs Integration

python

Google Cloud Healthcare API Integration

class GoogleHealthcareAPIClient:

"""Native integration with Google Cloud Healthcare APIs"""

def __init__(self, project_id: str, location: str, dataset_id: str):

self.project_id = project_id

self.location = location

self.dataset_id = dataset_id

self.healthcare_client = healthcare.HealthcareServiceClient()

self.fhir_client = self.setup_fhir_client()

self.hl7v2_client = self.setup_hl7v2_client()

self.dicom_client = self.setup_dicom_client()

self.consent_client = self.setup_consent_client()

def setup_fhir_client(self) -> FHIRStoreClient:

"""Setup FHIR R4 store client for structured healthcare data"""

fhir_store_path = self.healthcare_client.fhir_store_path(

self.project_id, self.location, self.dataset_id, "sarathi-fhir-store"

)

return FHIRStoreClient(fhir_store_path)

async def store_patient_data(self, patient_resource: dict) -> str:

"""Store patient data in FHIR format"""

try:

Validate FHIR resource

validated_resource = **await** self.validate_fhir_resource(

patient_resource, "Patient"

)

Store in Google Cloud Healthcare FHIR store

response = **await** self.fhir_client.create_resource(

parent=self.fhir_store_path,

body=validated_resource

)

Log to Cloud Audit Logs

await self.log_healthcare_data_access(

operation="create_patient",

resource_id=response.name,

data_classification="PHI"

)

return response.name


```
except Exception as e:
```

```
    await self.handle_healthcare_api_error(e, "store_patient_data")
```

```
    raise
```

```
async def retrieve_patient_consent(self, patient_id: str) -> ConsentRecord:
```

```
    """Retrieve patient consent using Google Cloud Healthcare Consent API"""
```

```
    try:
```

```
        consent_store_path = self.healthcare_client.consent_store_path(
            self.project_id, self.location, self.dataset_id, "sarathi-consent-store"
        )
```

```
        # Query consent records
```

```
        consent_response = await self.consent_client.list_consents(
            parent=consent_store_path,
            filter=f'user_id="{patient_id}"'
        )
```

```
        return ConsentRecord.from_google_response(consent_response)
```

```
except Exception as e:
```

```
    await self.handle_healthcare_api_error(e, "retrieve_patient_consent")
```

```
    raise
```

```
async def process_hl7v2_message(self, hl7_message: str) -> ProcessedMessage:
```

```
    """Process HL7v2 messages for legacy system integration"""
```

```
    try:
```

```
        hl7v2_store_path = self.healthcare_client.hl7v2_store_path(
            self.project_id, self.location, self.dataset_id, "sarathi-hl7v2-store"
        )
```

```
        # Ingest HL7v2 message
```

```
        message_response = await self.hl7v2_client.ingest_message(
            parent=hl7v2_store_path,
            message={"data": base64.b64encode(hl7_message.encode()).decode()}
        )
```

```
        # Extract structured data
```

```
        parsed_data = await self.parse_hl7v2_message(message_response)
```

```
        return ProcessedMessage(
```

```
            message_id=message_response.name,
```

```
            parsed_data=parsed_data,
```

```
            processing_timestamp=datetime.utcnow())
```

)

except Exception as e:

await self.handle_healthcare_api_error(e, "process_hl7v2_message")

raise

Google Cloud Healthcare Data Pipeline

class HealthcareDataPipeline:

"""Healthcare data processing pipeline using Google Cloud services"""

def __init__(self):

self.healthcare_client = GoogleHealthcareAPIClient(

project_id="sarathi-healthcare-platform",

location="us-central1",

dataset_id="sarathi-production-dataset"

)

self.nlp_client = HealthcareNLPClient()

self.autoML_client = AutoMLClient()

async def process_clinical_document(self, document: ClinicalDocument) -> ProcessedClinicalData:

"""Process clinical documents using Google Healthcare NLP API"""

Extract medical entities using Healthcare NLP API

nlp_response = await self.nlp_client.analyze_entities(

document_content=document.content,

license_type="HEALTHCARE"

)

Extract structured data

medical_entities = self.extract_medical_entities(nlp_response)

Store in FHIR format

fhir_resources = await self.convert_to_fhir_resources(medical_entities)

stored_resources = []

for resource in fhir_resources:

resource_id = await self.healthcare_client.store_fhir_resource(resource)

stored_resources.append(resource_id)

return ProcessedClinicalData(

document_id=document.id,

medical_entities=medical_entities,

fhir_resources=stored_resources,

processing_timestamp=datetime.utcnow())

)

class HealthcareNLPClient:

"""Google Cloud Healthcare NLP API client"""

def __init__(self):

self.nlp_service = language.LanguageServiceClient()

async def analyze_entities(self, document_content: str,

license_type: str = "HEALTHCARE") -> dict:

"""Analyze medical entities in clinical text"""

Configure healthcare-specific NLP

document = language.Document(

content=document_content,

type_=language.Document.Type.PLAIN_TEXT,

language="en"

)

Use healthcare-licensed NLP

features = language.AnnotateTextRequest.Features(

extract_entities=True,

extract_entity_sentiment=False,

extract_syntax=False,

classify_text=False,

extract_document_sentiment=False

)

response = await self.nlp_service.annotate_text(

request={

"document": document,

"features": features,

"encoding_type": language.EncodingType.UTF8

}

)

Extract healthcare-specific entities

medical_entities = self.extract_healthcare_entities(response.entities)

return {

"entities": medical_entities,

"confidence_scores": self.calculate_confidence_scores(response.entities),

"processing_metadata": {

"api_version": "healthcare_nlp_v1",

```
        "license_type": license_type,  
        "language": "en"  
    }  
}
```

2. Google AI (Gemini) Integration

```
python
```

```
# Google AI (Gemini) Integration for Healthcare
```

```
class GoogleAIGeminiClient:
```

```
    """Native integration with Google AI Gemini models for healthcare"""
```

```
    def __init__(self):
```

```
        self.vertex_ai_client = aiplatform.gapic.PredictionServiceClient()
```

```
        self.project_id = "sarathi-healthcare-platform"
```

```
        self.location = "us-central1"
```

```
        self.gemini_endpoint = self.setup_gemini_endpoint()
```

```
    def setup_gemini_endpoint(self) -> str:
```

```
        """Setup Gemini Pro endpoint for healthcare applications"""
```

```
        return f"projects/{self.project_id}/locations/{self.location}/publishers/google/models/gemini-pro-healthcare"
```

```
    async def execute_healthcare_agent(self, agent_config: HealthcareAgentConfig,
```

```
        input_data: dict) -> GeminiResponse:
```

```
        """Execute healthcare AI agent using Gemini Pro"""
```

```
        try:
```

```
            # Prepare healthcare-specific prompt
```

```
            healthcare_prompt = await self.prepare_healthcare_prompt(
```

```
                agent_config, input_data
```

```
            )
```

```
            # Configure Gemini for healthcare use
```

```
            gemini_request = {
```

```
                "endpoint": self.gemini_endpoint,
```

```
                "instances": [{
```

```
                    "prompt": healthcare_prompt,
```

```
                    "parameters": {
```

```
                        "temperature": agent_config.temperature,
```

```
                        "max_output_tokens": agent_config.max_tokens,
```

```
                        "top_p": agent_config.top_p,
```

```
                        "top_k": agent_config.top_k,
```

```
                        "safety_settings": self.get_healthcare_safety_settings(),
```

```
                        "healthcare_mode": True, # Enable healthcare-specific features
```

```
                        "phi_protection": True, # Enable PHI protection
```

```
                        "medical_accuracy": "high" # Prioritize medical accuracy
```

```
                    }
```

```
                ]
```

```
            }
```

```
            # Execute with healthcare monitoring
```

```

start_time = time.time()
response = await self.vertex_ai_client.predict(
    request=gemini_request
)
execution_time = (time.time() - start_time) * 1000

# Validate healthcare response
validated_response = await self.validate_healthcare_response(
    response, agent_config
)

# Log healthcare AI usage
await self.log_healthcare_ai_usage(
    agent_config=agent_config,
    input_data=input_data,
    response=validated_response,
    execution_time=execution_time
)

return GeminiResponse(
    content=validated_response.predictions[0]["content"],
    safety_ratings=validated_response.predictions[0]["safety_ratings"],
    execution_time=execution_time,
    model_version="gemini-pro-healthcare",
    healthcare_validated=True,
    phi_detected=validated_response.phi_analysis.phi_detected,
    medical_accuracy_score=validated_response.medical_accuracy_score
)

except Exception as e:
    await self.handle_gemini_error(e, agent_config)
    raise

def get_healthcare_safety_settings(self) -> list:
    """Get healthcare-specific safety settings for Gemini"""
    return [
        {
            "category": "HARM_CATEGORY_MEDICAL_ADVICE",
            "threshold": "BLOCK_MEDIUM_AND_ABOVE"
        },
        {
            "category": "HARM_CATEGORY_HEALTH_CLAIMS",
            "threshold": "BLOCK_LOW_AND_ABOVE"
        },
    ],

```

```
{
    "category": "HARM_CATEGORY_PHI_EXPOSURE",
    "threshold": "BLOCK_NONE" # We handle PHI separately
},
{
    "category": "HARM_CATEGORY_MEDICAL_MISINFORMATION",
    "threshold": "BLOCK_LOW_AND_ABOVE"
}
]
```

```
async def prepare_healthcare_prompt(self, agent_config: HealthcareAgentConfig,
                                     input_data: dict) -> str:
```

```
    """Prepare healthcare-optimized prompt for Gemini"""
```

```
    base_prompt = f"""
```

You are a Google AI assistant specialized in healthcare applications for the Sarthi platform.

Healthcare Context:

- Agent Type: {agent_config.agent_type}
- Medical Specialty: {agent_config.medical_specialty}
- Compliance Level: HIPAA, GDPR healthcare
- Safety Level: Maximum medical safety protocols

Instructions:

```
{agent_config.instructions}
```

Medical Guidelines:

- Follow evidence-based medical practices
- Cite relevant medical literature when applicable
- Maintain patient confidentiality and privacy
- Never provide definitive diagnoses
- Always recommend healthcare provider consultation
- Use appropriate medical terminology

Input Data:

```
{json.dumps(input_data, indent=2)}
```

Please provide a response following these healthcare standards and the specific agent instructions.

```
"""
```

```
return base_prompt
```

```
async def validate_healthcare_response(self, response: any,
                                       agent_config: HealthcareAgentConfig) -> ValidatedResponse:
```

```
"""Validate Gemini response for healthcare compliance"""
```

```
validation_result = ValidatedResponse()
```

```
# Extract response content
```

```
response_content = response.predictions[0]["content"]
```

```
# PHI detection and protection
```

```
phi_analysis = await self.analyze_phi_in_response(response_content)
```

```
validation_result.phi_analysis = phi_analysis
```

```
if phi_analysis.phi_detected and not agent_config.allow_phi_output:
```

```
# Redact PHI from response
```

```
response_content = await self.redact_phi_from_response(  
    response_content, phi_analysis.phi_locations  
)
```

```
# Medical accuracy validation
```

```
medical_accuracy_score = await self.validate_medical_accuracy(  
    response_content, agent_config.medical_specialty  
)
```

```
validation_result.medical_accuracy_score = medical_accuracy_score
```

```
# Safety validation
```

```
safety_validation = await self.validate_healthcare_safety(response_content)  
validation_result.safety_validation = safety_validation
```

```
# Update response with validated content
```

```
response.predictions[0]["content"] = response_content  
validation_result.predictions = response.predictions
```

```
return validation_result
```

```
# Google AI Model Manager for Healthcare
```

```
class GoogleAIModelManager:
```

```
    """Manage Google AI models for healthcare applications"""
```

```
    def __init__(self):
```

```
        self.vertex_ai = aiplatform
```

```
        self.model_registry = VertexAIModelRegistry()
```

```
        self.auto_ml_client = AutoMLClient()
```

```
    async def deploy_custom_healthcare_model(self, model_config: HealthcareModelConfig) -> str:
```

```
        """Deploy custom healthcare model using Vertex AI"""
```


Create custom healthcare model

```
model = aiplatform.Model.upload(
    display_name=model_config.model_name,
    artifact_uri=model_config.model_artifact_uri,
    serving_container_image_uri=model_config.container_image_uri,
    serving_container_health_route="/health",
    serving_container_predict_route="/predict",
    labels={
        "healthcare": "true",
        "hipaa_compliant": "true",
        "medical_specialty": model_config.medical_specialty,
        "sarathi_agent": model_config.agent_type
    }
)
```

Deploy with healthcare-specific configuration

```
endpoint = model.deploy(
    endpoint_display_name=f"{model_config.model_name}-endpoint",
    machine_type="n1-standard-4",
    min_replica_count=2,
    max_replica_count=10,
    accelerator_type="NVIDIA_TESLA_T4",
    accelerator_count=1,
    encryption_spec_key_name=model_config.kms_key_name, # Healthcare encryption
    enable_request_response_logging=True, # For compliance
    request_response_logging_sampling_rate=1.0 # Log all for audit
)
```

Configure healthcare monitoring

```
await self.setup_healthcare_model_monitoring(endpoint, model_config)
```

```
return endpoint.resource_name
```

```
async def setup_healthcare_model_monitoring(self, endpoint: aiplatform.Endpoint,
                                           model_config: HealthcareModelConfig):
```

```
    """Setup comprehensive monitoring for healthcare models"""
```

Create model monitoring job

```
monitoring_job = aiplatform.ModelDeploymentMonitoringJob.create(
    display_name=f"{model_config.model_name}-monitoring",
    endpoint=endpoint,
    logging_sampling_strategy=aiplatform.SamplingStrategy(
        random_sample_config=aiplatform.RandomSampleConfig(
```

```

        sample_rate=1.0 # Monitor all requests for healthcare
    )
),
model_deployment_monitoring_objective_configs=[
    aiplatform.ModelDeploymentMonitoringObjectiveConfig(
        deployed_model_id=endpoint.list_models()[0].id,
        objective_config=aiplatform.ModelMonitoringObjectiveConfig(
            training_dataset=aiplatform.InputDataConfig(
                dataset=model_config.training_dataset_uri
            ),
            training_prediction_skew_detection_config=aiplatform.TrainingPredictionSkewDetectionConfig(
                skew_thresholds={
                    "medical_accuracy": 0.05, # Low tolerance for medical accuracy drift
                    "confidence_score": 0.1,
                    "response_quality": 0.05
                }
            )
        )
    ],
    model_deployment_monitoring_schedule_config=aiplatform.ModelDeploymentMonitoringScheduleConfig(
        monitor_interval=3600 # Check every hour for healthcare models
    )
)

```

```

return monitoring_job

```

Google Cloud Healthcare Event Processing

```

class HealthcareEventProcessor:

```

```

    """Process healthcare events using Google Cloud Pub/Sub and Eventarc"""

```

```

    def __init__(self):

```

```

        self.publisher_client = pubsub_v1.PublisherClient()

```

```

        self.subscriber_client = pubsub_v1.SubscriberClient()

```

```

        self.healthcare_topic = f"projects/{PROJECT_ID}/topics/sarathi-healthcare-events"

```

```

    async def publish_healthcare_event(self, event: HealthcareEvent) -> str:

```

```

        """Publish healthcare event to Google Cloud Pub/Sub"""

```

Prepare healthcare event message

```

    event_message = {

```

```

        "event_id": str(uuid.uuid4()),

```

```

        "event_type": event.event_type,

```

```

        "timestamp": datetime.utcnow().isoformat(),

```

```
"patient_id": event.patient_id,
"facility_id": event.facility_id,
"data_classification": event.data_classification,
"phi_level": event.phi_level,
"event_data": event.data
}
```

```
# Encrypt PHI if present
```

```
if event.phi_level > 0:
    event_message["event_data"] = await self.encrypt_healthcare_data(
        event.data, event.kms_key_name
    )
```

```
# Publish with healthcare attributes
```

```
message_data = json.dumps(event_message).encode("utf-8")
future = self.publisher_client.publish(
    self.healthcare_topic,
    message_data,
    event_type=event.event_type,
    phi_level=str(event.phi_level),
    facility_id=event.facility_id,
    data_classification=event.data_classification
)
```

```
message_id = await future.result()
```

```
# Log healthcare event publication
```

```
await self.log_healthcare_event_publication(event, message_id)
```

```
return message_id
```

```
async def process_healthcare_event(self, message: pubsub_v1.PubsubMessage) -> ProcessingResult:
```

```
    """Process incoming healthcare event"""
```

```
try:
```

```
    # Parse healthcare event
```

```
    event_data = json.loads(message.data.decode("utf-8"))
```

```
    # Validate healthcare event structure
```

```
    validated_event = await self.validate_healthcare_event(event_data)
```

```
    # Decrypt PHI if present
```

```
    if validated_event.phi_level > 0:
        decrypted_data = await self.decrypt_healthcare_data(
```

```
        validated_event.event_data, validated_event.kms_key_name
    )
    validated_event.event_data = decrypted_data

    # Route to appropriate healthcare processor
    processor_result = await self.route_healthcare_event(validated_event)

    # Log successful processing
    await self.log_healthcare_event_processing(validated_event, processor_result)

    return ProcessingResult(
        success=True,
        event_id=validated_event.event_id,
        processing_time=processor_result.processing_time,
        output_data=processor_result.output_data
    )

except Exception as e:
    await self.handle_healthcare_event_error(e, message)
    raise
```

3. Google Cloud Technology Stack

yaml

Google Cloud Healthcare Technology Stack

google_cloud_services:

compute_and_containers:

primary: "Google Kubernetes Engine (GKE Autopilot)"
serverless: "Cloud Run (Healthcare compliant)"
functions: "Cloud Functions 2nd Gen"
batch_processing: "Cloud Batch"

ai_and_ml:

primary_llm: "Gemini Pro (Healthcare)"
secondary_llm: "Gemini Pro Vision (Medical Imaging)"
custom_models: "Vertex AI Custom Models"
automl: "Vertex AI AutoML (Healthcare)"
nlp: "Healthcare Natural Language API"
document_ai: "Document AI (Healthcare)"

healthcare_apis:

fhir_store: "Cloud Healthcare FHIR Store (R4/STU3)"
dicom_store: "Cloud Healthcare DICOM Store"
hl7v2_store: "Cloud Healthcare HL7v2 Store"
consent_management: "Cloud Healthcare Consent API"
de_identification: "Cloud Healthcare De-identification API"

databases_and_storage:

primary_db: "Cloud SQL for PostgreSQL (Healthcare)"
document_db: "Firestore (Healthcare mode)"
analytics_db: "BigQuery (Healthcare)"
global_db: "Cloud Spanner (Multi-region)"
cache: "Memorystore for Redis (HA)"
time_series: "Cloud Bigtable"
object_storage: "Cloud Storage (Healthcare)"

messaging_and_events:

pub_sub: "Cloud Pub/Sub (Healthcare events)"
workflows: "Cloud Workflows (Healthcare processes)"
tasks: "Cloud Tasks (Background jobs)"
eventarc: "Eventarc (Event-driven architecture)"
scheduler: "Cloud Scheduler"

security_and_compliance:

identity: "Cloud Identity and Access Management (IAM)"
key_management: "Cloud Key Management Service (KMS)"
secret_management: "Secret Manager"

vpc_security: "VPC Service Controls"
armor: "Cloud Armor (WAF)"
binary_authorization: "Binary Authorization"
audit_logs: "Cloud Audit Logs"
dlp: "Cloud Data Loss Prevention API"

monitoring_and_observability:
 monitoring: "Cloud Monitoring (Healthcare dashboards)"
 logging: "Cloud Logging (Healthcare audit)"
 tracing: "Cloud Trace"
 profiler: "Cloud Profiler"
 error_reporting: "Error Reporting"

networking:
 load_balancer: "Cloud Load Balancing"
 cdn: "Cloud CDN"
 dns: "Cloud DNS"
 interconnect: "Cloud Interconnect (Healthcare partners)"

development_and_deployment:
 build: "Cloud Build (Healthcare CI/CD)"
 deploy: "Cloud Deploy"
 artifact_registry: "Artifact Registry"
 source_repos: "Cloud Source Repositories"

Google Cloud Healthcare Configuration

google_healthcare_config:
 project_id: "sarathi-healthcare-platform"
 location: "us-central1" # HIPAA compliant region

healthcare_dataset:
 name: "sarathi-production-dataset"
 time_zone: "UTC"

fhir_stores:
 - name: "sarathi-fhir-r4-store"
 version: "R4"
 enable_update_create: true
 disable_referential_integrity: false

dicom_stores:
 - name: "sarathi-dicom-store"
 notification_config:
 pubsub_topic: "projects/sarathi-healthcare-platform/topics/dicom-events"

hl7v2_stores:

- name: "sarathi-hl7v2-store"

parser_config:

allow_null_header: false

segment_terminator: "\\r"

consent_stores:

- name: "sarathi-consent-store"

enable_consent_create_on_update: true

default_consent_ttl: "315360000s" # 10 years

security_config:

kms_key_name: "projects/sarathi-healthcare-platform/locations/us-central1/keyRings/sarathi-healthcare/cryptoKeys/sa"

audit_log_config:

log_type: "ADMIN_READ"

exempted_members: []

4. Google Cloud Deployment Configuration

yaml

Google Cloud Healthcare Deployment

apiVersion: apps/v1

kind: Deployment

metadata:

name: sadp-api

namespace: sarthi-healthcare

labels:

app: sadp-api

healthcare: "true"

hipaa-compliant: "true"

spec:

replicas: 3

selector:

matchLabels:

app: sadp-api

template:

metadata:

labels:

app: sadp-api

healthcare: "true"

annotations:

Google Cloud specific annotations

run.googleapis.com/vpc-access-connector: "sarthi-vpc-connector"

run.googleapis.com/vpc-access-egress: "private-ranges-only"

spec:

serviceAccountName: sadp-healthcare-service-account

securityContext:

runAsNonRoot: true

runAsUser: 65534

fsGroup: 65534

containers:

- name: sadp-api

image: gcr.io/sarthi-healthcare-platform/sadp-api:v1.0.0

ports:

- containerPort: 8080

name: http

env:

Google Cloud Healthcare API configuration

- name: GOOGLE_CLOUD_PROJECT

value: "sarthi-healthcare-platform"

- name: HEALTHCARE_DATASET_ID

value: "sarthi-production-dataset"

- name: HEALTHCARE_LOCATION

value: "us-central1"

Gemini AI configuration

- name: GEMINI_MODEL_ENDPOINT

value: "projects/sarathi-healthcare-platform/locations/us-central1/publishers/google/models/gemini-pro-healthcare"

- name: VERTEX_AI_PROJECT

value: "sarathi-healthcare-platform"

- name: VERTEX_AI_LOCATION

value: "us-central1"

Google Cloud services configuration

- name: CLOUD_SQL_CONNECTION_NAME

valueFrom:

secretKeyRef:

name: sarathi-db-secret

key: connection-name

- name: REDIS_HOST

valueFrom:

secretKeyRef:

name: sarathi-redis-secret

key: host

- name: PUBSUB_TOPIC

value: "projects/sarathi-healthcare-platform/topics/sarathi-healthcare-events"

Security and compliance

- name: KMS_KEY_NAME

value: "projects/sarathi-healthcare-platform/locations/us-central1/keyRings/sarathi-healthcare/cryptoKeys/sarathi-pl"

- name: HEALTHCARE_COMPLIANCE_MODE

value: "HIPAA"

resources:

requests:

memory: "1Gi"

cpu: "500m"

limits:

memory: "2Gi"

cpu: "1000m"

livenessProbe:

httpGet:

path: /health

port: 8080

initialDelaySeconds: 30

periodSeconds: 10

readinessProbe:

httpGet:

path: /ready

port: 8080

initialDelaySeconds: 5

periodSeconds: 5

Google Cloud security

securityContext:

allowPrivilegeEscalation: false

capabilities:

drop:

- ALL

readOnlyRootFilesystem: true

Google Cloud specific service configuration

apiVersion: v1

kind: Service

metadata:

name: sadp-api-service

namespace: sarthi-healthcare

annotations:

cloud.google.com/backend-config: '{"default": "sadp-backend-config"}'

cloud.google.com/load-balancer-type: "External"

spec:

type: LoadBalancer

loadBalancerSourceRanges:

- "10.0.0.0/8" *# Internal VPC only*

selector:

app: sadp-api

ports:

- port: 80

targetPort: 8080

protocol: TCP

Google Cloud Backend Configuration

apiVersion: cloud.google.com/v1

kind: BackendConfig

metadata:

name: sadp-backend-config

namespace: sarthi-healthcare

```
spec:
  healthCheck:
    checkIntervalSec: 10
    timeoutSec: 5
    healthyThreshold: 2
    unhealthyThreshold: 3
    type: HTTP
    requestPath: /health
  sessionAffinity:
    affinityType: "CLIENT_IP"
  timeoutSec: 300
  connectionDraining:
    drainingTimeoutSec: 60
  # Healthcare-specific security
  securityPolicy:
    name: "sarathi-healthcare-security-policy"
  iap:
    enabled: true
    oauthclientCredentials:
      secretName: "oauth-client-secret"

---
# Google Cloud IAM Service Account
apiVersion: v1
kind: ServiceAccount
metadata:
  name: sadp-healthcare-service-account
  namespace: sarathi-healthcare
  annotations:
    iam.gke.io/gcp-service-account: sadp-healthcare@sarathi-healthcare-platform.iam.gserviceaccount.com
```

4. Advanced API Specifications with OpenAPI 3.0

```
yaml
```

openapi: 3.0.3

info:

title: Sarthi AI Agent Development Platform API

version: 1.0.0

description: |

Enterprise-grade API for AI agent management, evaluation, and execution.

Supports healthcare-specific workflows with HIPAA compliance.

contact:

name: Sarthi Platform Team

url: <https://sarthi.com/support>

email: support@sarthi.com

license:

name: Commercial License

url: <https://sarthi.com/license>

servers:

- url: <https://sadb.sarthi.com/api/v1>

description: Production server

- url: <https://sadb-staging.sarthi.com/api/v1>

description: Staging server

security:

- BearerAuth: []

- ApiKeyAuth: []

paths:

/agents/{agentName}/execute:

post:

summary: Execute AI agent

description: |

Execute a specific AI agent with input data and configuration.

Supports healthcare-specific agents with HIPAA compliance.

operationId: executeAgent

tags:

- Agent Runtime

parameters:

- name: agentName

in: path

required: true

schema:

type: string

enum: [

"document_processor", "clinical_agent", "billing_agent",

"voice_agent", "health_assistant", "medication_entry",
"referral_processing", "lab_result_entry"

]

example: "clinical_agent"

- name: X-Trace-ID

in: header

schema:

type: string

format: uuid

description: Unique trace ID for request tracking

- name: X-Facility-ID

in: header

schema:

type: string

description: Healthcare facility identifier

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/AgentExecutionRequest'

examples:

clinical_agent_example:

summary: Clinical agent execution

value:

input_data:

patient_id: "SARTHI-PT-001"

primary_diagnosis: "Type 2 Diabetes Mellitus"

patient_age: 55

comorbidities: ["Hypertension", "Obesity"]

context:

facility_id: "SARTHI-CLINIC-001"

provider_id: "DR-SMITH-001"

session_id: "sess_123456"

options:

timeout: 30000

priority: "high"

trace_enabled: true

responses:

'200':

description: Agent execution successful

content:

application/json:

schema:

\$ref: '#/components/schemas/AgentExecutionResponse'

'400':

\$ref: '#/components/responses/BadRequest'

'401':

\$ref: '#/components/responses/Unauthorized'

'429':

\$ref: '#/components/responses/RateLimited'

'500':

\$ref: '#/components/responses/InternalServerError'

/workflows/execute:

post:

summary: Execute multi-agent workflow

description: |

Execute a complex workflow involving multiple AI agents
with orchestration and dependency management.

operationId: executeWorkflow

tags:

- Workflow Orchestration

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/WorkflowExecutionRequest'

responses:

'200':

description: Workflow execution initiated

content:

application/json:

schema:

\$ref: '#/components/schemas/WorkflowExecutionResponse'

'202':

description: Workflow accepted for asynchronous processing

content:

application/json:

schema:

\$ref: '#/components/schemas/AsyncWorkflowResponse'

/evaluation/agents/{agentName}/test:

post:

summary: Run comprehensive agent evaluation

description: |

Execute comprehensive testing and evaluation of an AI agent

against predefined test suites and benchmarks.

operationId: evaluateAgent

tags:

- Evaluation & Testing

parameters:

- name: agentName

in: path

required: true

schema:

type: string

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/AgentEvaluationRequest'

responses:

'200':

description: Evaluation completed

content:

application/json:

schema:

\$ref: '#/components/schemas/AgentEvaluationResponse'

'202':

description: Evaluation initiated asynchronously

content:

application/json:

schema:

\$ref: '#/components/schemas/AsyncEvaluationResponse'

/monitoring/agents/{agentName}/metrics:

get:

summary: Get agent performance metrics

description: |

Retrieve real-time and historical performance metrics
for a specific AI agent.

operationId: getAgentMetrics

tags:

- Monitoring & Analytics

parameters:

- name: agentName

in: path

required: true

schema:

type: string
- name: timeRange
in: query
schema:
type: string
enum: ["1h", "24h", "7d", "30d"]
default: "24h"
- name: metrics
in: query
schema:
type: array
items:
type: string
enum: ["accuracy", "latency", "error_rate", "cost", "throughput"]
style: form
explode: false
- name: aggregation
in: query
schema:
type: string
enum: ["mean", "median", "p95", "p99", "sum"]
default: "mean"

responses:

'200':

description: Metrics retrieved successfully

content:

application/json:

schema:

\$ref: '#/components/schemas/AgentMetricsResponse'

components:

securitySchemes:

BearerAuth:

type: http

scheme: bearer

bearerFormat: JWT

ApiKeyAuth:

type: apiKey

in: header

name: X-API-Key

schemas:

AgentExecutionRequest:

type: object

required:

- input_data

properties:

input_data:

type: object

description: Agent-specific input data

additionalProperties: true

context:

\$ref: '#/components/schemas/ExecutionContext'

options:

\$ref: '#/components/schemas/ExecutionOptions'

ExecutionContext:

type: object

properties:

facility_id:

type: string

description: Healthcare facility identifier

provider_id:

type: string

description: Healthcare provider identifier

patient_id:

type: string

description: Patient identifier (if applicable)

session_id:

type: string

description: Session identifier for tracking

compliance_level:

type: string

enum: ["HIPAA", "GDPR", "STANDARD"]

default: "HIPAA"

ExecutionOptions:

type: object

properties:

timeout:

type: integer

minimum: 1000

maximum: 300000

default: 30000

description: Timeout in milliseconds

priority:

type: string

enum: ["low", "normal", "high", "urgent"]

default: "normal"

trace_enabled:

type: boolean

default: false

description: Enable detailed execution tracing

cache_enabled:

type: boolean

default: true

description: Enable response caching

fallback_enabled:

type: boolean

default: true

description: Enable fallback mechanisms

AgentExecutionResponse:

type: object

required:

- execution_id
- agent_name
- status
- result

properties:

execution_id:

type: string

format: uuid

description: Unique execution identifier

agent_name:

type: string

description: Name of the executed agent

status:

type: string

enum: ["completed", "failed", "timeout", "cancelled"]

execution_time:

type: integer

description: Execution time in milliseconds

result:

type: object

description: Agent execution result

additionalProperties: true

metadata:

\$ref: '#/components/schemas/ExecutionMetadata'

compliance:

\$ref: '#/components/schemas/ComplianceInfo'

performance:

\$ref: '#/components/schemas/PerformanceMetrics'

error:

\$ref: '#/components/schemas/ErrorInfo'

ExecutionMetadata:

type: object

properties:

prompt_version:

type: string

description: Version of the prompt used

model_info:

type: object

properties:

provider:

type: string

example: "anthropic"

model:

type: string

example: "claude-sonnet-4-20250514"

version:

type: string

tokens_used:

type: object

properties:

input:

type: integer

output:

type: integer

total:

type: integer

cost:

type: number

format: decimal

description: Execution cost in USD

cache_hit:

type: boolean

description: Whether response was served from cache

ComplianceInfo:

type: object

properties:

hipaa_compliant:

type: boolean

description: HIPAA compliance status

audit_trail_id:
type: string
description: Audit trail identifier
phi_processed:
type: boolean
description: Whether PHI was processed
encryption_used:
type: boolean
description: Whether data was encrypted
retention_policy:
type: string
description: Data retention policy applied

PerformanceMetrics:

type: object
properties:
accuracy_score:
type: number
minimum: 0
maximum: 100
description: Accuracy score (0-100)
confidence_score:
type: number
minimum: 0
maximum: 1
description: Confidence score (0-1)
latency_percentile:
type: string
enum: ["p50", "p90", "p95", "p99"]
description: Latency percentile classification
quality_score:
type: number
minimum: 0
maximum: 100
description: Overall quality score

ErrorInfo:

type: object
properties:
error_code:
type: string
description: Error code for categorization
error_message:
type: string

description: Human-readable error message

error_details:

type: object

description: Detailed error information

additionalProperties: true

retry_possible:

type: boolean

description: Whether the operation can be retried

suggested_action:

type: string

description: Suggested action to resolve the error

responses:

BadRequest:

description: Invalid request parameters

content:

application/json:

schema:

\$ref: '#/components/schemas/ErrorResponse'

Unauthorized:

description: Authentication required

content:

application/json:

schema:

\$ref: '#/components/schemas/ErrorResponse'

RateLimited:

description: Rate limit exceeded

content:

application/json:

schema:

\$ref: '#/components/schemas/ErrorResponse'

headers:

Retry-After:

schema:

type: integer

description: Seconds to wait before retrying

InternalServerError:

description: Internal server error

content:

application/json:

schema:

`$ref: '#/components/schemas/ErrorResponse'`

ErrorResponse:

type: object

required:

- error
- message
- timestamp

properties:

error:

type: string

description: Error type

message:

type: string

description: Error message

details:

type: object

description: Additional error details

additionalProperties: true

timestamp:

type: string

format: date-time

description: Error timestamp

trace_id:

type: string

description: Request trace ID for debugging

support_id:

type: string

description: Support ticket ID for complex issues

5. Google Cloud Security & Compliance Framework

python

```
# Google Cloud Healthcare Security Implementation
```

```
class GoogleCloudHealthcareSecurity:
```

```
    """Google Cloud native security for healthcare applications"""
```

```
    def __init__(self):
```

```
        self.iam_client = iam.IamPolicyManagementServiceClient()
```

```
        self.kms_client = kms.KeyManagementServiceClient()
```

```
        self.dlp_client = dlp.DlpServiceClient()
```

```
        self.audit_logger = GoogleCloudAuditLogger()
```

```
        self.vpc_sc_client = accesscontextmanager.AccessContextManagerClient()
```

```
    async def authenticate_with_google_iam(self, request: Request) -> GoogleServiceContext:
```

```
        """Google Cloud IAM-based authentication for healthcare APIs"""
```

```
        # Extract Google Cloud Identity token
```

```
        id_token = self.extract_google_identity_token(request)
```

```
        # Verify token with Google Cloud IAM
```

```
        try:
```

```
            # Validate JWT token
```

```
            claims = id_token.verify_oauth2_token(
```

```
                id_token,
```

```
                google.auth.transport.requests.Request(),
```

```
                audience=self.get_expected_audience()
```

```
            )
```

```
            # Get service account details
```

```
            service_account = await self.get_service_account_details(
```

```
                claims["email"]
```

```
            )
```

```
            # Validate healthcare permissions
```

```
            healthcare_roles = await self.get_healthcare_iam_roles(service_account)
```

```
        return GoogleServiceContext(
```

```
            service_account_email=claims["email"],
```

```
            project_id=claims["aud"].split("/")[1],
```

```
            iam_roles=healthcare_roles,
```

```
            healthcare_permissions=self.extract_healthcare_permissions(healthcare_roles),
```

```
            authenticated_at=datetime.utcnow()
```

```
        )
```

```
    except ValueError as e:
```

```
await self.audit_logger.log_authentication_failure(
    error="invalid_token",
    request_details=self.sanitize_request_for_logging(request)
)
raise InvalidAuthenticationError("Invalid Google Cloud identity token")
```

```
async def authorize_healthcare_action(self, context: GoogleServiceContext,
                                     resource: str, action: str) -> bool:
    """Google Cloud IAM-based authorization for healthcare operations"""
```

```
# Check Google Cloud IAM permissions
```

```
required_permission = f"healthcare.{resource}.{action}"
```

```
try:
```

```
# Test IAM permissions
```

```
permissions_response = await self.iam_client.test_iam_permissions(
    resource=f"projects/{context.project_id}/locations/us-central1/datasets/sarhi-production-dataset",
    permissions=[required_permission]
)
```

```
has_permission = required_permission in permissions_response.permissions
```

```
if not has_permission:
```

```
    await self.audit_logger.log_authorization_failure(
        service_account=context.service_account_email,
        resource=resource,
        action=action,
        reason="insufficient_iam_permissions"
    )
```

```
    return False
```

```
# Additional healthcare-specific checks
```

```
if not await self.validate_healthcare_context(context, resource, action):
    return False
```

```
return True
```

```
except Exception as e:
```

```
    await self.audit_logger.log_authorization_error(e, context)
    return False
```

```
async def encrypt_phi_with_google_kms(self, phi_data: dict,
                                     patient_id: str) -> EncryptedData:
    """Encrypt PHI using Google Cloud KMS with healthcare keys"""
```



```

# Get healthcare KMS key
key_name = f"projects/{self.project_id}/locations/us-central1/keyRings/sarathi-healthcare/cryptoKeys/sarathi-phi-key"

# Detect PHI fields
phi_fields = await self.detect_phi_with_google_dlp(phi_data)

encrypted_data = phi_data.copy()
encryption_metadata = []

for field_path, field_value in phi_fields.items():
    # Encrypt each PHI field separately
    plaintext = json.dumps(field_value).encode('utf-8')

    # Additional authenticated data for context
    additional_data = f"patient_id:{patient_id},field:{field_path}".encode('utf-8')

    # Encrypt with Google Cloud KMS
    encrypt_response = await self.kms_client.encrypt(
        request={
            "name": key_name,
            "plaintext": plaintext,
            "additional_authenticated_data": additional_data
        }
    )

    # Store encrypted data
    encrypted_data[field_path] = {
        "encrypted_value": base64.b64encode(encrypt_response.ciphertext).decode('utf-8'),
        "key_version": encrypt_response.name,
        "encryption_algorithm": "GOOGLE_SYMMETRIC_ENCRYPTION"
    }

    encryption_metadata.append({
        "field_path": field_path,
        "key_version": encrypt_response.name,
        "patient_id": patient_id,
        "encrypted_at": datetime.utcnow().isoformat()
    })

# Log encryption for audit
await self.audit_logger.log_phi_encryption(
    patient_id=patient_id,
    fields_encrypted=list(phi_fields.keys()),

```

```
key_version=key_name,  
encryption_metadata=encryption_metadata  
)
```

```
return EncryptedData(  
    encrypted_data=encrypted_data,  
    encryption_metadata=encryption_metadata,  
    phi_fields_count=len(phi_fields)  
)
```

```
async def detect_phi_with_google_dlp(self, data: dict) -> dict:  
    """Detect PHI using Google Cloud Data Loss Prevention API"""
```

```
# Configure DLP for healthcare PHI detection
```

```
inspect_config = {  
    "info_types": [  
        {"name": "PERSON_NAME"},  
        {"name": "PHONE_NUMBER"},  
        {"name": "EMAIL_ADDRESS"},  
        {"name": "US_SOCIAL_SECURITY_NUMBER"},  
        {"name": "DATE_OF_BIRTH"},  
        {"name": "MEDICAL_RECORD_NUMBER"},  
        {"name": "US_HEALTHCARE_NPI"},  
        {"name": "CREDIT_CARD_NUMBER"},  
        {"name": "US_BANK_ACCOUNT"},  
        {"name": "IBAN_CODE"},  
        # Healthcare-specific info types  
        {"name": "US_DEA_NUMBER"},  
        {"name": "ICD9_CODE"},  
        {"name": "ICD10_CODE"},  
        {"name": "US_PASSPORT"},  
        {"name": "US_DRIVERS_LICENSE_NUMBER"}  
    ],  
    "custom_info_types": [  
        {  
            "info_type": {"name": "PATIENT_ID"},  
            "regex": {  
                "pattern": r"SARTHI-PT-\d{6}"  
            }  
        }  
    ],  
    "min_likelihood": "POSSIBLE",  
    "limits": {  
        "max_findings_per_info_type": 100
```

```

    },
    "include_quote": True
}

# Prepare data for inspection
content_item = {
    "value": json.dumps(data)
}

# Call Google Cloud DLP
dlp_response = await self.dlp_client.inspect_content(
    request={
        "parent": f"projects/{self.project_id}/locations/global",
        "inspect_config": inspect_config,
        "item": content_item
    }
)

# Extract PHI findings
phi_fields = {}
for finding in dlp_response.result.findings:
    field_path = self.extract_field_path_from_quote(finding.quote, data)
    phi_fields[field_path] = {
        "info_type": finding.info_type.name,
        "likelihood": finding.likelihood.name,
        "quote": finding.quote,
        "location": finding.location
    }

return phi_fields

async def apply_google_vpc_service_controls(self, request: Request,
                                           context: GoogleServiceContext) -> bool:
    """Apply Google Cloud VPC Service Controls for healthcare data"""

    # Check if request is within authorized VPC Service Controls perimeter
    perimeter_name = f"accessPolicies/{self.access_policy_id}/servicePerimeters/sarathi-healthcare-perimeter"

    try:
        # Get current service perimeter
        perimeter = await self.vpc_sc_client.get_service_perimeter(
            name=perimeter_name
        )

```

```
# Validate request context against perimeter
```

```
request_context = {  
    "origin": {  
        "ip_address": request.client.host,  
        "user_agent": request.headers.get("user-agent")  
    },  
    "destination": {  
        "service": "healthcare.googleapis.com",  
        "method": request.method,  
        "resource": request.url.path  
    }  
}
```

```
# Check perimeter restrictions
```

```
if not self.validate_perimeter_access(perimeter, request_context, context):  
    await self.audit_logger.log_vpc_service_control_violation(  
        perimeter_name=perimeter_name,  
        request_context=request_context,  
        service_context=context  
    )  
    return False
```

```
    return True
```

```
except Exception as e:  
    await self.audit_logger.log_vpc_service_control_error(e, context)  
    return False
```

```
class GoogleCloudHealthcareCompliance:
```

```
    """Google Cloud native compliance for healthcare"""
```

```
    def __init__(self):  
        self.audit_log_client = logging.Client()  
        self.healthcare_client = healthcare.HealthcareServiceClient()  
        self.bigquery_client = bigquery.Client()  
        self.dlp_client = dlp.DlpServiceClient()
```

```
    async def log_healthcare_api_access(self, access_log: HealthcareAccessLog):  
        """Log healthcare API access to Google Cloud Audit Logs"""
```

```
# Prepare structured audit log entry
```

```
audit_entry = {  
    "protoPayload": {  
        "@type": "type.googleapis.com/google.cloud.audit.AuditLog",
```

```
"serviceName": "sarathi-sadp.googleapis.com",
"methodName": f"/{access_log.api_version}/{access_log.method}",
"resourceName": access_log.resource_name,
"authenticationInfo": {
    "principalEmail": access_log.service_account_email,
    "serviceAccountKeyName": access_log.service_account_key_name
},
"authorizationInfo": [
    {
        "resource": access_log.resource_name,
        "permission": access_log.required_permission,
        "granted": access_log.permission_granted
    }
],
"requestMetadata": {
    "callerIp": access_log.caller_ip,
    "callerSuppliedUserAgent": access_log.user_agent,
    "requestAttributes": {
        "time": access_log.timestamp.isoformat(),
        "reason": "HEALTHCARE_API_ACCESS"
    }
},
"request": {
    "healthcare_operation": access_log.operation_type,
    "phi_accessed": access_log.phi_accessed,
    "patient_count": access_log.patient_count,
    "data_classification": access_log.data_classification
},
"response": {
    "status": access_log.response_status,
    "processing_time_ms": access_log.processing_time
},
},
"insertId": str(uuid.uuid4()),
"resource": {
    "type": "healthcare_dataset",
    "labels": {
        "project_id": self.project_id,
        "location": "us-central1",
        "dataset_id": "sarathi-production-dataset"
    }
},
"timestamp": access_log.timestamp.isoformat(),
"severity": "INFO",
```

```

        "labels": {
            "healthcare_compliance": "HIPAA",
            "phi_level": str(access_log.phi_level),
            "facility_id": access_log.facility_id
        }
    }
}

```

Write to Google Cloud Audit Logs

```
self.audit_log_client.write_entries([audit_entry])
```

Also store in BigQuery for analytics

```
await self.store_compliance_data_in_bigquery(access_log)
```

async def generate_hipaa_audit_report(self, report_config: HIPAAAuditReportConfig) -> AuditReport:

```
    """Generate HIPAA audit report using Google Cloud services"""
```

Query audit logs from BigQuery

```
audit_query = f"""
```

```
SELECT
```

```
    timestamp,
```

```
    protoPayload.authenticationInfo.principalEmail as service_account,
```

```
    protoPayload.resourceName as resource,
```

```
    protoPayload.request.healthcare_operation as operation,
```

```
    protoPayload.request.phi_accessed as phi_accessed,
```

```
    protoPayload.request.patient_count as patient_count,
```

```
    protoPayload.response.status as status,
```

```
    labels.facility_id as facility_id
```

```
FROM `self.project_id.sarthi_audit_logs.healthcare_access_logs`
```

```
WHERE timestamp >= @start_date
```

```
AND timestamp <= @end_date
```

```
AND labels.healthcare_compliance = 'HIPAA'
```

```
ORDER BY timestamp DESC
```

```
"""
```

```
query_job = self.bigquery_client.query(
```

```
    audit_query,
```

```
    job_config=bigquery.QueryJobConfig(
```

```
        query_parameters=[
```

```
            bigquery.ScalarQueryParameter("start_date", "TIMESTAMP", report_config.start_date),
```

```
            bigquery.ScalarQueryParameter("end_date", "TIMESTAMP", report_config.end_date)
```

```
        ]
```

```
    )
```

```
)
```

```
audit_results = query_job.result()
```

```
# Analyze audit data
```

```
report_data = {  
    "report_period": {  
        "start_date": report_config.start_date.isoformat(),  
        "end_date": report_config.end_date.isoformat()  
    },  
    "total_api_calls": 0,  
    "phi_access_events": 0,  
    "unique_patients_accessed": set(),  
    "service_accounts_used": set(),  
    "facilities_accessed": set(),  
    "compliance_violations": [],  
    "access_patterns": {}  
}
```

```
for row in audit_results:
```

```
    report_data["total_api_calls"] += 1
```

```
    if row.phi_accessed:
```

```
        report_data["phi_access_events"] += 1
```

```
        report_data["unique_patients_accessed"].add(row.patient_count)
```

```
    report_data["service_accounts_used"].add(row.service_account)
```

```
    report_data["facilities_accessed"].add(row.facility_id)
```

```
# Check for compliance violations
```

```
if await self.detect_compliance_violation(row):
```

```
    report_data["compliance_violations"].append({  
        "timestamp": row.timestamp,  
        "violation_type": "UNAUTHORIZED_PHI_ACCESS",  
        "service_account": row.service_account,  
        "resource": row.resource  
    })
```

```
# Convert sets to counts for JSON serialization
```

```
report_data["unique_patients_count"] = len(report_data["unique_patients_accessed"])
```

```
report_data["unique_service_accounts"] = len(report_data["service_accounts_used"])
```

```
report_data["unique_facilities"] = len(report_data["facilities_accessed"])
```

```
# Remove sets (not JSON serializable)
```

```
del report_data["unique_patients_accessed"]
```

```
del report_data["service_accounts_used"]
```

```
del report_data["facilities_accessed"]
```

```
return AuditReport(  
    report_id=str(uuid.uuid4()),  
    generated_at=datetime.utcnow(),  
    report_type="HIPAA_COMPLIANCE",  
    data=report_data,  
    compliance_status="COMPLIANT" if not report_data["compliance_violations"] else "VIOLATIONS_DETECTED"  
)
```

```
# Google Cloud Healthcare API Key Management
```

```
class GoogleCloudAPIKeyManager:
```

```
    """Google Cloud native API key management for healthcare services"""
```

```
def __init__(self):
```

```
    self.api_keys_client = apikeys_v2.ApiKeysClient()
```

```
    self.iam_client = iam.IamPolicyManagementServiceClient()
```

```
    self.secret_manager_client = secretmanager.SecretManagerServiceClient()
```

```
async def create_healthcare_api_key(self, key_config: HealthcareAPIKeyConfig) -> GoogleHealthcareAPIKey:
```

```
    """Create Google Cloud API key for healthcare services"""
```

```
# Create API key with healthcare-specific restrictions
```

```
api_key_request = {
```

```
    "parent": f"projects/{self.project_id}/locations/global",
```

```
    "api_key": {
```

```
        "display_name": key_config.display_name,
```

```
        "restrictions": {
```

```
            "api_targets": [
```

```
                {
```

```
                    "service": "healthcare.googleapis.com",
```

```
                    "methods": key_config.allowed_methods
```

```
                },
```

```
                {
```

```
                    "service": "aiplatform.googleapis.com",
```

```
                    "methods": ["predict", "explain"]
```

```
                }
```

```
            ],
```

```
            "server_key_restrictions": {
```

```
                "allowed_ips": key_config.allowed_ip_ranges
```

```
            }
```

```
        },
```

```
        "annotations": {
```

```
            "healthcare_compliant": "true",
```



```

        "facility_id": key_config.facility_id,
        "service_type": key_config.service_type,
        "phi_access_level": str(key_config.phi_access_level)
    }
}

# Create the API key
operation = await self.api_keys_client.create_key(request=api_key_request)
api_key_response = await operation.result()

# Store key metadata in Secret Manager
secret_data = {
    "api_key_id": api_key_response.name,
    "key_string": api_key_response.key_string,
    "created_at": datetime.utcnow().isoformat(),
    "facility_id": key_config.facility_id,
    "phi_access_level": key_config.phi_access_level
}

secret_name = f"projects/{self.project_id}/secrets/sadp-api-key-{key_config.service_type}-{key_config.facility_id}"

await self.secret_manager_client.create_secret(
    request={
        "parent": f"projects/{self.project_id}",
        "secret_id": f"sadp-api-key-{key_config.service_type}-{key_config.facility_id}",
        "secret": {
            "replication": {
                "user_managed": {
                    "replicas": [
                        {"location": "us-central1"},
                        {"location": "us-east1"} # Backup region
                    ]
                }
            },
            "labels": {
                "healthcare": "true",
                "facility_id": key_config.facility_id,
                "service_type": key_config.service_type
            }
        }
    }
)

```

```
# Add secret version
```

```
await self.secret_manager_client.add_secret_version(  
    request={  
        "parent": secret_name,  
        "payload": {"data": json.dumps(secret_data).encode("utf-8")}  
    }  
)
```

```
return GoogleHealthcareAPIKey(  
    api_key_id=api_key_response.name,  
    key_string=api_key_response.key_string,  
    restrictions=api_key_response.restrictions,  
    facility_id=key_config.facility_id,  
    phi_access_level=key_config.phi_access_level,  
    created_at=datetime.utcnow()  
)
```

```
async def validate_healthcare_api_key(self, api_key_string: str) -> Optional[HealthcareAPIKeyContext]:  
    """Validate Google Cloud API key for healthcare operations"""
```

```
try:
```

```
    # Look up API key
```

```
    lookup_response = await self.api_keys_client.lookup_key(  
        request={"key_string": api_key_string}  
    )
```

```
    api_key = lookup_response.parent
```

```
    # Get key details
```

```
    key_details = await self.api_keys_client.get_key(name=api_key)
```

```
    # Validate key is active and not expired
```

```
    if key_details.state != apikeys_v2.Key.State.ACTIVE:  
        return None
```

```
    # Extract healthcare context from annotations
```

```
    annotations = key_details.annotations
```

```
return HealthcareAPIKeyContext(  
    api_key_id=key_details.name,  
    facility_id=annotations.get("facility_id"),  
    service_type=annotations.get("service_type"),  
    phi_access_level=int(annotations.get("phi_access_level", "0")),  
    healthcare_compliant=annotations.get("healthcare_compliant") == "true",
```

```
        allowed_services=self.extract_allowed_services(key_details.restrictions),
        key_created_at=key_details.create_time
    )
```

```
except Exception as e:
```

```
    # Log validation attempt
```

```
    await self.audit_logger.log_api_key_validation_error(e, api_key_string)
```

```
    return None
```

```
# Updated Google Cloud Security Configuration
```

```
GOOGLE_CLOUD_SECURITY_CONFIG = {
```

```
    "authentication": {
```

```
        "method": "google_cloud_iam",
```

```
        "service_account_auth": True,
```

```
        "identity_token_validation": True,
```

```
        "api_key_validation": True
```

```
    },
```

```
    "authorization": {
```

```
        "iam_permissions": True,
```

```
        "healthcare_role_validation": True,
```

```
        "vpc_service_controls": True,
```

```
        "resource_level_permissions": True
```

```
    },
```

```
    "encryption": {
```

```
        "kms_provider": "google_cloud_kms",
```

```
        "key_ring": "sarathi-healthcare",
```

```
        "phi_encryption_key": "sarathi-phi-key",
```

```
        "field_level_encryption": True,
```

```
        "encryption_at_rest": True,
```

```
        "encryption_in_transit": True
```

```
    },
```

```
    "compliance": {
```

```
        "phi_detection": "google_cloud_dlp",
```

```
        "audit_logging": "google_cloud_audit_logs",
```

```
        "data_residency": "us_central1",
```

```
        "vpc_service_controls": True,
```

```
        "binary_authorization": True
```

```
    },
```

```
    "monitoring": {
```

```
        "security_monitoring": "google_cloud_security_command_center",
```

```
        "audit_analysis": "google_cloud_bigquery",
```

```
        "threat_detection": "google_cloud_armor",
```

```
        "anomaly_detection": "google_cloud_monitoring"
```

```
    }
```

}

```## Google Cloud Healthcare API Specifications

### Enhanced API Specifications for Google Cloud Healthcare

```yaml

openapi: 3.0.3

info:

title: Sarthi AI Agent Development Platform API - Google Healthcare

version: 1.0.0

description: |

Enterprise-grade API for AI agent management using Google Cloud Healthcare APIs and Google AI (Gemini) models. Supports healthcare-specific workflows with Google Cloud native HIPAA compliance.

contact:

name: Sarthi Platform Team

url: <https://sarthi.com/support>

email: support@sarthi.com

license:

name: Commercial License

url: <https://sarthi.com/license>

servers:

- url: <https://sadp.sarthi.com/api/v1>

description: Production server (Google Cloud)

- url: <https://sadp-staging.sarthi.com/api/v1>

description: Staging server (Google Cloud)

security:

- GoogleCloudAuth: []

- GoogleAPIKey: []

paths:

[/agents/{agentName}/execute](#):

post:

summary: Execute AI agent with Google AI (Gemini)

description: |

Execute a healthcare AI agent using Google Cloud Healthcare APIs and Google AI Gemini models with HIPAA compliance.

operationId: executeHealthcareAgent

tags:

- Google Healthcare Agents

parameters:

- name: agentName

in: path
required: true
schema:
 type: string
 enum: [
 "gemini_clinical_agent", "gemini_billing_agent",
 "healthcare_nlp_agent", "fhir_processing_agent",
 "medical_imaging_agent", "hl7_processing_agent"
]

example: "gemini_clinical_agent"

- name: X-Google-Healthcare-Dataset

in: header
required: true
schema:
 type: string
description: Google Cloud Healthcare dataset ID
example: "projects/sarathi-healthcare-platform/locations/us-central1/datasets/sarathi-production-dataset"

- name: X-Google-Project-ID

in: header
required: true
schema:
 type: string
description: Google Cloud Project ID
example: "sarathi-healthcare-platform"

requestBody:

required: true
content:
 application/json:
 schema:
 \$ref: '#/components/schemas/GoogleHealthcareAgentRequest'

examples:

 gemini_clinical_example:

 summary: Gemini clinical agent execution

 value:

 input_data:

 patient_fhir_resource:

 resourceType: "Patient"

 id: "SARTHI-PT-001"

 name: [{"family": "Doe", "given": ["Jane"]}]

 clinical_context:

 primary_diagnosis: "Type 2 Diabetes Mellitus"

 icd10_code: "E11.9"

 google_healthcare_context:

 fhir_store: "projects/sarathi-healthcare-platform/locations/us-central1/datasets/sarathi-production-dataset/"

```
consent_store: "projects/sarathi-healthcare-platform/locations/us-central1/datasets/sarathi-production-data"
gemini_config:
  model: "gemini-pro-healthcare"
  temperature: 0.2
  max_tokens: 2048
  safety_settings: "healthcare_maximum"
google_cloud_options:
  use_healthcare_nlp: true
  phi_detection_enabled: true
  audit_logging_level: "FULL"
```

responses:

'200':

description: Healthcare agent execution successful

content:

application/json:

schema:

\$ref: '#/components/schemas/GoogleHealthcareAgentResponse'

'400':

\$ref: '#/components/responses/BadRequest'

'401':

\$ref: '#/components/responses/Unauthorized'

'403':

\$ref: '#/components/responses/GoogleCloudPermissionDenied'

/healthcare/fhir/process:

post:

summary: Process FHIR resources with Google Healthcare APIs

description: |

Process FHIR resources using Google Cloud Healthcare FHIR stores
with integrated AI analysis via Gemini models.

operationId: processFHIRResources

tags:

- Google Healthcare FHIR

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/FHIRProcessingRequest'

responses:

'200':

description: FHIR processing completed

content:

application/json:

schema:

\$ref: '#/components/schemas/FHIRProcessingResponse'

/healthcare/hl7v2/process:

post:

summary: Process HL7v2 messages with Google Healthcare APIs

description: |

Process HL7v2 messages using Google Cloud Healthcare HL7v2 stores with legacy system integration.

operationId: processHL7v2Messages

tags:

- Google Healthcare HL7v2

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/HL7v2ProcessingRequest'

responses:

'200':

description: HL7v2 processing completed

content:

application/json:

schema:

\$ref: '#/components/schemas/HL7v2ProcessingResponse'

/healthcare/consent/validate:

post:

summary: Validate patient consent with Google Healthcare Consent API

description: |

Validate patient consent for data access using Google Cloud Healthcare Consent API with privacy controls.

operationId: validatePatientConsent

tags:

- Google Healthcare Consent

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/ConsentValidationRequest'

responses:

'200':

description: Consent validation completed

content:

application/json:

schema:

\$ref: '#/components/schemas/ConsentValidationResponse'

/google-ai/gemini/healthcare:

post:

summary: Direct Gemini healthcare model access

description: |

Direct access to Google AI Gemini models configured for healthcare applications with enhanced safety settings.

operationId: geminiHealthcareInference

tags:

- Google AI Gemini

requestBody:

required: true

content:

application/json:

schema:

\$ref: '#/components/schemas/GeminiHealthcareRequest'

responses:

'200':

description: Gemini inference completed

content:

application/json:

schema:

\$ref: '#/components/schemas/GeminiHealthcareResponse'

components:

securitySchemes:

GoogleCloudAuth:

type: oauth2

flows:

clientCredentials:

tokenUrl: https://oauth2.googleapis.com/token

scopes:

https://www.googleapis.com/auth/cloud-healthcare: "Access Google Cloud Healthcare APIs"

https://www.googleapis.com/auth/cloud-platform: "Access Google Cloud Platform"

GoogleAPIKey:

type: apiKey

in: header

name: X-Goog-Api-Key

schemas:

GoogleHealthcareAgentRequest:

type: object

required:

- input_data
- gemini_config

properties:

input_data:

type: object

description: Healthcare-specific input data

properties:

patient_fhir_resource:

\$ref: '#/components/schemas/FHIRResource'

clinical_context:

type: object

properties:

primary_diagnosis:

type: string

example: "Type 2 Diabetes Mellitus"

icd10_code:

type: string

example: "E11.9"

symptoms:

type: array

items:

type: string

google_healthcare_context:

\$ref: '#/components/schemas/GoogleHealthcareContext'

gemini_config:

\$ref: '#/components/schemas/GeminiConfig'

google_cloud_options:

\$ref: '#/components/schemas/GoogleCloudOptions'

GoogleHealthcareContext:

type: object

properties:

fhir_store:

type: string

description: Google Cloud Healthcare FHIR store path

example: "projects/sarhi-healthcare-platform/locations/us-central1/datasets/sarhi-production-dataset/fhirStore"

dicom_store:

type: string

description: Google Cloud Healthcare DICOM store path

hl7v2_store:

type: string

description: Google Cloud Healthcare HL7v2 store path

consent_store:

type: string

description: Google Cloud Healthcare Consent store path

dataset_id:

type: string

example: "sarhi-production-dataset"

location:

type: string

example: "us-central1"

GeminiConfig:

type: object

required:

- model

properties:

model:

type: string

enum: ["gemini-pro-healthcare", "gemini-pro-vision-healthcare"]

default: "gemini-pro-healthcare"

temperature:

type: number

minimum: 0.0

maximum: 2.0

default: 0.2

description: Lower temperature for healthcare accuracy

max_tokens:

type: integer

minimum: 1

maximum: 8192

default: 2048

top_p:

type: number

minimum: 0.0

maximum: 1.0

default: 0.8

top_k:

type: integer

minimum: 1

maximum: 40

default: 40

safety_settings:

type: string

enum: ["healthcare_maximum", "healthcare_standard", "standard"]

default: "healthcare_maximum"

description: Healthcare-specific safety configurations

healthcare_mode:

type: boolean

default: true

description: Enable healthcare-specific features

phi_protection:

type: boolean

default: true

description: Enable PHI protection mechanisms

GoogleCloudOptions:

type: object

properties:

use_healthcare_nlp:

type: boolean

default: true

description: Use Google Healthcare NLP API for entity extraction

phi_detection_enabled:

type: boolean

default: true

description: Enable Google Cloud DLP for PHI detection

audit_logging_level:

type: string

enum: ["NONE", "BASIC", "FULL"]

default: "FULL"

description: Google Cloud Audit Logs level

kms_encryption:

type: boolean

default: true

description: Use Google Cloud KMS for encryption

vpc_service_controls:

type: boolean

default: true

description: Apply VPC Service Controls

GoogleHealthcareAgentResponse:

type: object

required:

- execution_id
- agent_name
- status
- result

properties:

execution_id:

type: string

format: uuid

description: Unique execution identifier

agent_name:

type: string

description: Name of the executed healthcare agent

status:

type: string

enum: ["completed", "failed", "timeout", "cancelled"]

execution_time:

type: integer

description: Execution time in milliseconds

result:

type: object

description: Healthcare agent execution result

properties:

clinical_analysis:

type: string

description: Clinical analysis from Gemini

confidence_score:

type: number

minimum: 0

maximum: 1

medical_entities:

type: array

items:

\$ref: '#/components/schemas/MedicalEntity'

recommendations:

type: array

items:

type: string

google_cloud_metadata:

\$ref: '#/components/schemas/GoogleCloudMetadata'

gemini_metadata:

\$ref: '#/components/schemas/GeminiMetadata'

compliance:

\$ref: '#/components/schemas/GoogleHealthcareCompliance'

performance:

\$ref: '#/components/schemas/PerformanceMetrics'

GoogleCloudMetadata:

type: object

properties:

project_id:

type: string

example: "sarathi-healthcare-platform"

region:

type: string

example: "us-central1"

healthcare_dataset:

type: string

example: "sarathi-production-dataset"

services_used:

type: array

items:

type: string

example: ["healthcare.googleapis.com", "aiplatform.googleapis.com"]

kms_key_used:

type: string

example: "projects/sarathi-healthcare-platform/locations/us-central1/keyRings/sarathi-healthcare/cryptoKeys/sarathi-healthcare-key"

GeminiMetadata:

type: object

properties:

model_version:

type: string

example: "gemini-pro-healthcare-001"

model_endpoint:

type: string

example: "us-central1-aiplatform.googleapis.com"

tokens_used:

type: object

properties:

input_tokens:

type: integer

output_tokens:

type: integer

total_tokens:

type: integer

safety_ratings:

type: array

items:

type: object

properties:

category:

type: string

probability:

type: string

blocked:

type: boolean

cost:

type: number

format: decimal

description: Execution cost in USD

GoogleHealthcareCompliance:

type: object

properties:

hipaa_compliant:

type: boolean

description: HIPAA compliance status

google_cloud_audit_logged:

type: boolean

description: Whether logged to Google Cloud Audit Logs

phi_detected:

type: boolean

description: Whether PHI was detected by Google Cloud DLP

phi_encrypted:

type: boolean

description: Whether PHI was encrypted with Google Cloud KMS

consent_validated:

type: boolean

description: Whether patient consent was validated

vpc_service_controls_applied:

type: boolean

description: Whether VPC Service Controls were applied

data_residency:

type: string

example: "us-central1"

description: Google Cloud data residency location

FHIRResource:

type: object

description: FHIR R4 resource structure

properties:

resourceType:

type: string

example: "Patient"

id:

type: string

example: "SARTHI-PT-001"

additionalProperties: true

MedicalEntity:

type: object

properties:

entity_type:

type: string

example: "MEDICAL_CONDITION"

text:

type: string

example: "Type 2 Diabetes"

confidence:

type: number

minimum: 0

maximum: 1

source:

type: string

enum: ["gemini", "healthcare_nlp", "manual"]

FHIRProcessingRequest:

type: object

required:

- fhir_resources
- processing_type

properties:

fhir_resources:

type: array

items:

\$ref: '#/components/schemas/FHIRResource'

processing_type:

type: string

enum: ["store", "search", "analyze", "validate"]

fhir_store_path:

type: string

description: Google Cloud Healthcare FHIR store path

analysis_config:

type: object

properties:

use_gemini_analysis:

type: boolean

default: true

extract_medical_entities:

type: boolean

default: true

FHIRProcessingResponse:

type: object

properties:

processing_id:

type: string

format: uuid

fhir_resources_processed:

type: integer

stored_resource_ids:

type: array

items:

type: string

analysis_results:

type: object

description: Results from Gemini analysis if requested

google_healthcare_metadata:

\$ref: '#/components/schemas/GoogleCloudMetadata'

HL7v2ProcessingRequest:

type: object

required:

- hl7_message
- processing_type

properties:

hl7_message:

type: string

description: Base64 encoded HL7v2 message

processing_type:

type: string

enum: ["ingest", "parse", "convert_to_fhir"]

hl7v2_store_path:

type: string

description: Google Cloud Healthcare HL7v2 store path

conversion_config:

type: object

properties:

target_fhir_version:

type: string

enum: ["R4", "STU3"]

default: "R4"

ConsentValidationRequest:

type: object

required:

- patient_id
- data_access_request

properties:

patient_id:

type: string

example: "SARTHI-PT-001"

data_access_request:

type: object

properties:

data_types:

type: array

items:

type: string

purpose:

type: string

example: "treatment"

requesting_organization:

type: string

consent_store_path:

type: string

description: Google Cloud Healthcare Consent store path

GeminiHealthcareRequest:

type: object

required:

- prompt
- gemini_config

properties:

prompt:

type: string

description: Healthcare-specific prompt for Gemini

gemini_config:

\$ref: '#/components/schemas/GeminiConfig'

context:

type: object

properties:

medical_specialty:

type: string

example: "cardiology"

patient_context:

type: object

description: Anonymized patient context

responses:

GoogleCloudPermissionDenied:

description: Google Cloud IAM permission denied

content:

application/json:

schema:

type: object

properties:

error:

type: string

example: "PERMISSION_DENIED"

message:

type: string

example: "Service account lacks required Google Cloud Healthcare permissions"

required_permissions:

type: array

items:

type: string

example: ["healthcare.datasets.get", "healthcare.fhirStores.search"]

google_cloud_project:

type: string

example: "sarathi-healthcare-platform"

Google Cloud Integration Examples

Complete Integration Example

python

```
# Complete Google Cloud Healthcare Integration Example
```

```
async def main():
```

```
    """Complete example of Google Cloud Healthcare Platform integration"""
```

```
    print("🏠 GOOGLE CLOUD HEALTHCARE PLATFORM INTEGRATION")
```

```
    print("=" * 60)
```

```
    # Initialize Google Cloud Healthcare components
```

```
    google_healthcare_client = GoogleHealthcareAPIClient(
```

```
        project_id="sarathi-healthcare-platform",
```

```
        location="us-central1",
```

```
        dataset_id="sarathi-production-dataset"
```

```
)
```

```
    gemini_client = GoogleAIGeminiClient()
```

```
    security_client = GoogleCloudHealthcareSecurity()
```

```
    compliance_client = GoogleCloudHealthcareCompliance()
```

```
    print("\n🔑 GOOGLE CLOUD AUTHENTICATION")
```

```
    print("-" * 40)
```

```
    # Authenticate with Google Cloud IAM
```

```
    mock_request = create_mock_request_with_google_auth()
```

```
    service_context = await security_client.authenticate_with_google_iam(mock_request)
```

```
    print(f"✅ Authenticated: {service_context.service_account_email}")
```

```
    print(f"   Project: {service_context.project_id}")
```

```
    print(f"   Healthcare Permissions: {len(service_context.healthcare_permissions)}")
```

```
    print("\n📄 2 FHIR RESOURCE PROCESSING")
```

```
    print("-" * 40)
```

```
    # Create sample FHIR patient resource
```

```
    patient_resource = {
```

```
        "resourceType": "Patient",
```

```
        "id": "SARTHI-PT-001",
```

```
        "name": [{"family": "Smith", "given": ["John"]}],
```

```
        "birthDate": "1970-05-15",
```

```
        "gender": "male",
```

```
        "identifier": [
```

```
            {
```

```
                "system": "http://sarathi.com/patient-id",
```

```
                "value": "SARTHI-PT-001"
```

```
            }
```

```
]
}
```

```
# Store in Google Cloud Healthcare FHIR store
```

```
stored_resource_id = await google_healthcare_client.store_patient_data(patient_resource)
print(f"✅ FHIR Patient stored: {stored_resource_id}")
```

```
print("\n 3 GEMINI HEALTHCARE AI PROCESSING")
print("-" * 40)
```

```
# Configure healthcare agent for Gemini
```

```
healthcare_agent_config = HealthcareAgentConfig(
    agent_type="gemini_clinical_agent",
    medical_specialty="primary_care",
    temperature=0.2,
    max_tokens=2048,
    allow_phi_output=False,
    safety_level="healthcare_maximum"
)
```

```
# Prepare clinical input
```

```
clinical_input = {
    "patient_fhir_resource": patient_resource,
    "chief_complaint": "Routine annual physical examination",
    "vital_signs": {
        "blood_pressure": "120/80",
        "heart_rate": 72,
        "temperature": "98.6°F",
        "weight": "180 lbs"
    },
    "medical_history": ["Hypertension", "Type 2 Diabetes"]
}
```

```
# Execute Gemini healthcare analysis
```

```
gemini_response = await gemini_client.execute_healthcare_agent(
    healthcare_agent_config, clinical_input
)
```

```
print(f"✅ Gemini Analysis Complete:")
print(f"  Model: {gemini_response.model_version}")
print(f"  Execution Time: {gemini_response.execution_time}ms")
print(f"  Healthcare Validated: {gemini_response.healthcare_validated}")
print(f"  Medical Accuracy Score: {gemini_response.medical_accuracy_score}")
print(f"  PHI Detected: {gemini_response.phi_detected}")
```

```
print("\n 4 GOOGLE CLOUD DLP PHI DETECTION")
```

```
print("-" * 40)
```

```
# Test PHI detection with Google Cloud DLP
```

```
test_data = {  
    "patient_name": "John Smith",  
    "ssn": "123-45-6789",  
    "medical_record_number": "MR-123456",  
    "diagnosis": "Type 2 Diabetes Mellitus"  
}
```

```
phi_fields = await security_client.detect_phi_with_google_dlp(test_data)
```

```
print(f"✅ PHI Detection Complete:")
```

```
print(f"  PHI Fields Found: {len(phi_fields)}")
```

```
for field, details in phi_fields.items():
```

```
    print(f"    - {field}: {details['info_type']} ({details['likelihood']})")
```

```
print("\n 5 GOOGLE CLOUD KMS ENCRYPTION")
```

```
print("-" * 40)
```

```
# Encrypt PHI using Google Cloud KMS
```

```
encrypted_data = await security_client.encrypt_phi_with_google_kms(  
    test_data, patient_id="SARTHI-PT-001"  
)
```

```
print(f"✅ PHI Encryption Complete:")
```

```
print(f"  Fields Encrypted: {encrypted_data.phi_fields_count}")
```

```
print(f"  Encryption Key: Google Cloud KMS")
```

```
print(f"  Key Ring: sarthi-healthcare")
```

```
print("\n 6 HEALTHCARE NLP PROCESSING")
```

```
print("-" * 40)
```

```
# Process clinical text with Healthcare NLP API
```

```
clinical_text = """
```

```
Patient presents with chest pain and shortness of breath.
```

```
History of hypertension and diabetes mellitus type 2.
```

```
Current medications include metformin 1000mg BID and lisinopril 10mg daily.
```

```
Vital signs: BP 140/90, HR 88, RR 18, Temp 98.4°F.
```

```
"""
```

```
nlp_client = HealthcareNLPClient()
```

```
nlp_results = await nlp_client.analyze_entities(clinical_text)
```

```

print(f"✅ Healthcare NLP Complete:")
print(f"  Medical Entities Found: {len(nlp_results['entities'])}")
print(f"  API Version: {nlp_results['processing_metadata']['api_version']}")

print("\n7 HL7V2 MESSAGE PROCESSING")
print("-" * 40)

# Sample HL7v2 message
hl7_message = """"MSH|^~\&|EPIC|SARTHI|||20250819120000||ADT^A01|123456|P|2.5
PID|1||SARTHI-PT-001^^^MR^MR||SMITH^JOHN^||19700515|M|||123 MAIN ST^^ANYTOWN^ST^12345^USA||(555)123-4567
PV1|1||2000^2012^01|||1234567890^ATTENDING^PHYSICIAN||SUR|||||1234567890^ATTENDING^PHYSICIAN|INP|A|||

processed_message = await google_healthcare_client.process_hl7v2_message(hl7_message)
print(f"✅ HL7v2 Processing Complete:")
print(f"  Message ID: {processed_message.message_id}")
print(f"  Processing Time: {processed_message.processing_timestamp}")

print("\n8 CONSENT VALIDATION")
print("-" * 40)

# Validate patient consent
consent_record = await google_healthcare_client.retrieve_patient_consent("SARTHI-PT-001")
print(f"✅ Consent Validation:")
print(f"  Consent Status: {consent_record.status if consent_record else 'Not Found'}")

print("\n9 COMPLIANCE AUDIT LOGGING")
print("-" * 40)

# Generate healthcare audit log
audit_log = HealthcareAccessLog(
    service_account_email=service_context.service_account_email,
    resource_name="projects/sarthi-healthcare-platform/locations/us-central1/datasets/sarthi-production-dataset",
    operation_type="FHIR_PATIENT_CREATE",
    phi_accessed=True,
    patient_count=1,
    data_classification="PHI",
    phi_level=3,
    facility_id="SARTHI-CLINIC-001"
)

await compliance_client.log_healthcare_api_access(audit_log)
print(f"✅ Audit Logging Complete:")
print(f"  Logged to Google Cloud Audit Logs")

```

```
print(f"  Stored in BigQuery for analytics")
```

```
print("\n📋 COMPREHENSIVE HIPAA AUDIT REPORT")
```

```
print("-" * 40)
```

```
# Generate HIPAA audit report
```

```
report_config = HIPAAAuditReportConfig(  
    start_date=datetime.utcnow() - timedelta(days=30),  
    end_date=datetime.utcnow(),  
    include_phi_access=True,  
    include_compliance_violations=True  
)
```

```
audit_report = await compliance_client.generate_hipaa_audit_report(report_config)
```

```
print(f"✅ HIPAA Audit Report Generated:")
```

```
print(f"  Report ID: {audit_report.report_id}")
```

```
print(f"  Compliance Status: {audit_report.compliance_status}")
```

```
print(f"  Total API Calls: {audit_report.data['total_api_calls']}")
```

```
print(f"  PHI Access Events: {audit_report.data['phi_access_events']}")
```

```
print("\n🌐 GOOGLE CLOUD HEALTHCARE INTEGRATION SUMMARY")
```

```
print("=" * 60)
```

```
print("✅ Google Cloud IAM authentication successful")
```

```
print("✅ Google Cloud Healthcare FHIR store integration")
```

```
print("✅ Google AI Gemini healthcare models")
```

```
print("✅ Google Cloud DLP PHI detection")
```

```
print("✅ Google Cloud KMS encryption")
```

```
print("✅ Healthcare NLP API integration")
```

```
print("✅ HL7v2 message processing")
```

```
print("✅ Patient consent validation")
```

```
print("✅ Google Cloud Audit Logs compliance")
```

```
print("✅ Comprehensive HIPAA audit reporting")
```

```
print(f"\n🏠 GOOGLE CLOUD NATIVE BENEFITS:")
```

```
print(f"  • Native HIPAA compliance with Google Cloud Healthcare APIs")
```

```
print(f"  • Advanced PHI protection with Google Cloud DLP and KMS")
```

```
print(f"  • Healthcare-optimized Gemini models")
```

```
print(f"  • Comprehensive audit trails with Cloud Audit Logs")
```

```
print(f"  • VPC Service Controls for data protection")
```

```
print(f"  • Google Cloud Security Command Center integration")
```

```
return {
```





```
    "google_cloud_integration": "successful",
```

```
    "gemini_healthcare_models": "operational",
```





```
"healthcare_apis": "integrated",  
"compliance_framework": "hipaa_compliant",  
"security_controls": "enterprise_grade"  
}  
  
if __name__ == "__main__":  
    # Run comprehensive Google Cloud Healthcare integration  
    asyncio.run(main())
```

Benefits of Google Cloud Healthcare Focus





1. Native Healthcare Compliance

-  **Built-in HIPAA compliance** with Google Cloud Healthcare APIs
-  **PHI protection** with Google Cloud DLP and KMS
-  **Audit trails** with Google Cloud Audit Logs
-  **Data residency** controls with Google Cloud regions





2. Advanced AI Capabilities

-  **Healthcare-optimized Gemini** models with medical safety
-  **Medical entity extraction** with Healthcare NLP API
-  **Medical imaging analysis** with Gemini Pro Vision
-  **Custom healthcare models** with Vertex AI





3. Comprehensive Healthcare APIs

-  **FHIR R4/STU3 support** with Cloud Healthcare FHIR stores
-  **HL7v2 integration** for legacy systems
-  **DICOM support** for medical imaging
-  **Consent management** with Healthcare Consent API

4. Enterprise Security

-  **Google Cloud IAM** for fine-grained access control
-  **VPC Service Controls** for data protection
-  **Binary Authorization** for container security
-  **Cloud Armor** for DDoS protection

5. Operational Excellence

-  **Managed services** reduce operational overhead
-  **Auto-scaling** with GKE Autopilot
-  **Monitoring** with Cloud Monitoring and Error Reporting
-  **Cost optimization** with Google Cloud pricing models

This Google Cloud Healthcare-focused architecture provides a robust, compliant, and scalable foundation for the Sarthi AI Agent Development Platform, leveraging Google's healthcare-specific infrastructure and AI capabilities.

6. Advanced Error Handling & Resilience

```
python
```

Circuit Breaker and Retry Patterns

class SADPResilienceFramework:

"""Advanced resilience patterns for SADP"""

def __init__(self):

self.circuit_breakers = {}

self.retry_policies = {

 "claude_api": ExponentialBackoff(

 initial_delay=1.0,

 max_delay=60.0,

 multiplier=2.0,

 max_retries=3

),

 "database": LinearBackoff(

 delay=0.5,

 max_retries=5

)

}

async def execute_with_resilience(self,

 operation_name: str,

 operation: Callable,

 *args, **kwargs) -> Any:

"""Execute operation with full resilience patterns"""

Circuit breaker check

circuit_breaker = self.get_circuit_breaker(operation_name)

if circuit_breaker.is_open():

 raise CircuitBreakerOpenError(f"Circuit breaker open for {operation_name}")

Retry with backoff

retry_policy = self.retry_policies.get(operation_name)

for attempt in range(retry_policy.max_retries + 1):

 try:

 result = await operation(*args, **kwargs)

 circuit_breaker.record_success()

 return result

 except RetryableError as e:

 circuit_breaker.record_failure()

 if attempt < retry_policy.max_retries:

```

        delay = retry_policy.calculate_delay(attempt)
        await asyncio.sleep(delay)
        continue
    else:
        raise e

```

```

except NonRetryableError as e:
    circuit_breaker.record_failure()
    raise e

```

```

def get_circuit_breaker(self, operation_name: str) -> CircuitBreaker:
    """Get or create circuit breaker for operation"""
    if operation_name not in self.circuit_breakers:
        self.circuit_breakers[operation_name] = CircuitBreaker(
            failure_threshold=5,
            timeout=30.0,
            expected_exception=Exception
        )
    return self.circuit_breakers[operation_name]

```

Advanced Error Classifications

```
class SADPErrorHandler:
```

```
    """Intelligent error handling and classification"""
```

```

ERROR_CLASSIFICATIONS = {
    "claude_api_errors": {
        "rate_limit": {"retryable": True, "severity": "medium"},
        "invalid_request": {"retryable": False, "severity": "low"},
        "server_error": {"retryable": True, "severity": "high"},
        "timeout": {"retryable": True, "severity": "medium"}
    },
    "validation_errors": {
        "schema_violation": {"retryable": False, "severity": "low"},
        "phi_leak_detected": {"retryable": False, "severity": "critical"},
        "compliance_violation": {"retryable": False, "severity": "critical"}
    },
    "infrastructure_errors": {
        "database_connection": {"retryable": True, "severity": "high"},
        "cache_miss": {"retryable": True, "severity": "low"},
        "service_unavailable": {"retryable": True, "severity": "high"}
    }
}

```

```
def classify_error(self, error: Exception) -> ErrorClassification:
```

```
"""Intelligent error classification"""
error_type = type(error).__name__
error_message = str(error)

# Use ML-based classification for unknown errors
if error_type not in self.ERROR_CLASSIFICATIONS:
    return self.ml_classify_error(error_message)

classification = self.ERROR_CLASSIFICATIONS[error_type]

return ErrorClassification(
    error_type=error_type,
    retryable=classification["retryable"],
    severity=classification["severity"],
    recommended_action=self.get_recommended_action(error),
    user_message=self.get_user_friendly_message(error)
)
```

7. Performance Optimization & Caching

python

Advanced Caching Strategy

class SADPCacheManager:

"""Multi-level caching with intelligent invalidation"""

def __init__(self):

self.l1_cache = TTLCache(maxsize=1000, ttl=300) *# 5 min in-memory*

self.l2_cache = redis.Redis(host="redis-cluster") *# 1 hour Redis*

self.l3_cache = self.setup_cdn_cache() *# 24 hour CDN*

async def get_cached_response(self, cache_key: str) -> Optional[dict]:

"""Multi-level cache retrieval"""

L1: In-memory cache (fastest)

if cache_key in self.l1_cache:

return self.l1_cache[cache_key]

L2: Redis cache (fast)

cached_data = await self.l2_cache.get(cache_key)

if cached_data:

data = json.loads(cached_data)

self.l1_cache[cache_key] = data *# Promote to L1*

return data

L3: CDN cache (for static responses)

if self.is_static_response(cache_key):

return await self.get_from_cdn(cache_key)

return None

async def cache_response(self, cache_key: str, data: dict, ttl: int = 3600):

"""Intelligent response caching"""

Determine cache levels based on data characteristics

cache_levels = self.determine_cache_levels(data)

if "l1" in cache_levels:

self.l1_cache[cache_key] = data

if "l2" in cache_levels:

await self.l2_cache.setex(

cache_key,

ttl,

json.dumps(data, cls=SADPJSONEncoder)

)

```
if "l3" in cache_levels and self.is_publicly_cacheable(data):  
    await self.cache_to_cdn(cache_key, data, ttl)
```

```
def determine_cache_levels(self, data: dict) -> List[str]:
```

```
    """Determine appropriate cache levels for data"""
```

```
    cache_levels = []
```

```
    # Always cache non-PHI data in L1
```

```
    if not self.contains_phi(data):  
        cache_levels.extend(["l1", "l2"])
```

```
    # Cache static responses in CDN
```

```
    if self.is_static_content(data):  
        cache_levels.append("l3")
```

```
    return cache_levels
```

```
# Connection Pooling and Resource Management
```

```
class SADPResourceManager:
```

```
    """Advanced resource management for external services"""
```

```
    def __init__(self):
```

```
        self.claude_client_pool = self.create_claude_pool()
```

```
        self.db_connection_pool = self.create_db_pool()
```

```
        self.redis_connection_pool = self.create_redis_pool()
```

```
    def create_claude_pool(self) -> ClientPool:
```

```
        """Create optimized Claude API client pool"""
```

```
        return ClientPool(  
            client_factory=lambda: anthropic.Anthropic(),  
            max_size=50,  
            min_size=10,  
            acquire_timeout=30.0,  
            max_lifetime=3600.0,  
            health_check_interval=60.0
```

```
)
```

```
    async def get_claude_client(self) -> anthropic.Anthropic:
```

```
        """Get Claude client with automatic retry and failover"""
```

```
        for attempt in range(3):
```

```
            try:
```

```
                client = await self.claude_client_pool.acquire()
```

```
# Health check
if await self.health_check_claude_client(client):
    return client
else:
    await self.claude_client_pool.discard(client)

except Exception as e:
    if attempt == 2: # Last attempt
        raise ClaudeClientUnavailableError(
            "Unable to acquire healthy Claude client"
        )
    await asyncio.sleep(2 ** attempt) # Exponential backoff
```

8. Advanced Monitoring & Observability

python

Comprehensive Monitoring Stack

class SADPMonitoringStack:

"""Production-grade monitoring and observability"""

def __init__(self):

self.prometheus_registry = CollectorRegistry()

self.metrics_collector = self.setup_metrics()

self.tracer = self.setup_distributed_tracing()

self.logger = self.setup_structured_logging()

def setup_metrics(self) -> MetricsCollector:

"""Setup comprehensive metrics collection"""

Business metrics

agent_execution_counter = Counter(

'sadp_agent_executions_total',

'Total number of agent executions',

['agent_name', 'status', 'facility_id']

)

agent_execution_duration = Histogram(

'sadp_agent_execution_duration_seconds',

'Agent execution duration in seconds',

['agent_name'],

buckets=(0.1, 0.5, 1.0, 2.5, 5.0, 10.0, 30.0, 60.0)

)

agent_accuracy_gauge = Gauge(

'sadp_agent_accuracy_score',

'Current agent accuracy score',

['agent_name']

)

Infrastructure metrics

claude_api_requests = Counter(

'sadp_claude_api_requests_total',

'Total Claude API requests',

['method', 'status_code']

)

claude_api_latency = Histogram(

'sadp_claude_api_latency_seconds',

'Claude API request latency',


```
    buckets=(0.1, 0.25, 0.5, 1.0, 2.5, 5.0, 10.0)
)
```

```
# Cost metrics
```

```
claude_token_usage = Counter(
    'sadp_claude_tokens_total',
    'Total Claude tokens consumed',
    ['agent_name', 'token_type']
)
```

```
daily_cost_gauge = Gauge(
    'sadp_daily_cost_usd',
    'Daily operational cost in USD',
    ['service']
)
```

```
return MetricsCollector(
    agent_execution_counter,
    agent_execution_duration,
    agent_accuracy_gauge,
    claude_api_requests,
    claude_api_latency,
    claude_token_usage,
    daily_cost_gauge
)
```

```
def setup_distributed_tracing(self) -> Tracer:
```

```
    """Setup distributed tracing with Jaeger"""
```

```
    config = Config(
        config={
            'sampler': {'type': 'const', 'param': 1},
            'logging': True,
            'reporter_batch_size': 1,
        },
        service_name='sadp-api',
        validate=True,
    )
```

```
    return config.initialize_tracer()
```

```
def setup_structured_logging(self) -> Logger:
```

```
    """Setup structured logging with correlation IDs"""
```

```
    logger = structlog.get_logger()
    structlog.configure(
        processors=[
```

```

        structlog.stdlib.filter_by_level,
        structlog.stdlib.add_logger_name,
        structlog.stdlib.add_log_level,
        structlog.processors.TimeStamper(fmt="iso"),
        structlog.processors.StackInfoRenderer(),
        structlog.processors.format_exc_info,
        structlog.processors.UnicodeDecoder(),
        structlog.processors.JSONRenderer()
    ],
    context_class=dict,
    logger_factory=structlog.stdlib.LoggerFactory(),
    wrapper_class=structlog.stdlib.BoundLogger,
    cache_logger_on_first_use=True,
)
return logger

```

Advanced Alerting System

class SADPAlertManager:

"""Intelligent alerting with context-aware notifications"""

def __init__(self):

```

    self.alert_rules = self.load_alert_rules()
    self.notification_channels = self.setup_notification_channels()
    self.escalation_policies = self.load_escalation_policies()

```

async def evaluate_alerts(self, metrics: dict):

"""Evaluate metrics against alert rules"""

for rule in self.alert_rules:

if await self.evaluate_rule(rule, metrics):

```

    alert = Alert(
        rule_name=rule.name,
        severity=rule.severity,
        message=rule.format_message(metrics),
        context=self.build_alert_context(metrics),
        timestamp=datetime.utcnow()
    )

```

await self.send_alert(alert)

async def send_alert(self, alert: Alert):

"""Send alert through appropriate channels"""

Determine notification channels based on severity

```
channels = self.get_channels_for_severity(alert.severity)
```

```
# Check for alert fatigue and deduplication
```

```
if not await self.should_send_alert(alert):  
    return
```

```
# Send notifications
```

```
tasks = []
```

```
for channel in channels:
```

```
    if channel.type == "slack":  
        tasks.append(self.send_slack_alert(alert, channel))  
    elif channel.type == "pagerduty":  
        tasks.append(self.send_pagerduty_alert(alert, channel))  
    elif channel.type == "email":  
        tasks.append(self.send_email_alert(alert, channel))
```

```
await asyncio.gather(*tasks, return_exceptions=True)
```

```
def load_alert_rules(self) -> List[AlertRule]:
```

```
    """Load alert rules from configuration"""
```

```
    return [
```

```
        AlertRule(  
            name="agent_accuracy_degradation",  
            condition="agent_accuracy < 90",  
            severity="warning",  
            cooldown=300 # 5 minutes
```

```
        ),
```

```
        AlertRule(  
            name="claude_api_high_latency",  
            condition="claude_api_p95_latency > 10",  
            severity="critical",  
            cooldown=60 # 1 minute
```

```
        ),
```

```
        AlertRule(  
            name="high_error_rate",  
            condition="error_rate > 0.05",  
            severity="critical",  
            cooldown=120 # 2 minutes
```

```
        ),
```

```
        AlertRule(  
            name="cost_spike",  
            condition="hourly_cost > daily_budget / 24 * 2",  
            severity="warning",  
            cooldown=1800 # 30 minutes
```

```
)  
]
```

9. Advanced Configuration Management

```
python
```

Environment-specific Configuration

class SADPConfig:

"""Comprehensive configuration management"""

def __init__(self, environment: str = "production"):

self.environment = environment

self.config = self.load_config()

self.secrets = self.load_secrets()

def load_config(self) -> dict:

"""Load environment-specific configuration"""

base_config = {

 "api": {

 "host": "0.0.0.0",

 "port": 8000,

 "workers": 4,

 "timeout": 300,

 "max_request_size": "50MB"

 },

 "database": {

 "pool_size": 20,

 "max_overflow": 30,

 "pool_timeout": 30,

 "pool_recycle": 3600

 },

 "redis": {

 "max_connections": 50,

 "socket_timeout": 5,

 "socket_connect_timeout": 5,

 "retry_on_timeout": True

 },

 "claude_api": {

 "timeout": 60,

 "max_retries": 3,

 "rate_limit": 100, # requests per minute

 "max_tokens": 4000

 },

 "security": {

 "jwt_expiration": 3600,

 "mfa_required_for": [

 "/api/v1/development/*",

 "/api/v1/evaluation/*/deploy"

],

```
"encryption_algorithm": "AES-256-GCM",
"password_policy": {
  "min_length": 12,
  "require_uppercase": True,
  "require_lowercase": True,
  "require_numbers": True,
  "require_symbols": True
},
"compliance": {
  "hipaa_enabled": True,
  "audit_retention_days": 2555, # 7 years
  "phi_encryption_required": True,
  "data_residency": "US",
  "backup_encryption": True
},
"monitoring": {
  "metrics_retention_days": 90,
  "log_retention_days": 365,
  "alert_cooldown_seconds": 300,
  "health_check_interval": 30
}
}
```

Environment-specific overrides

```
env_overrides = {
  "development": {
    "api": {"workers": 1, "debug": True},
    "database": {"pool_size": 5},
    "security": {"jwt_expiration": 86400}, # 24 hours
    "compliance": {"hipaa_enabled": False}
  },
  "staging": {
    "api": {"workers": 2},
    "database": {"pool_size": 10},
    "claude_api": {"rate_limit": 50}
  },
  "production": {
    "api": {"workers": 4, "debug": False},
    "database": {"pool_size": 20},
    "claude_api": {"rate_limit": 200}
  }
}
```

```
config = base_config.copy()
if self.environment in env_overrides:
    config.update(env_overrides[self.environment])

return config

def load_secrets(self) -> dict:
    """Load secrets from secure storage"""
    if self.environment == "development":
        return self.load_from_env()
    else:
        return self.load_from_vault()

def load_from_vault(self) -> dict:
    """Load secrets from HashiCorp Vault or Google Secret Manager"""
    # Implementation for production secret management
    pass
```

10. Deployment and DevOps Enhancements

yaml

Advanced Helm Chart Structure

sadp-platform/

- |— Chart.yaml
- |— values.yaml
- |— values-dev.yaml
- |— values-staging.yaml
- |— values-prod.yaml
- |— templates/
 - |— deployment.yaml
 - |— service.yaml
 - |— ingress.yaml
 - |— configmap.yaml
 - |— secrets.yaml
 - |— hpa.yaml
 - |— pdb.yaml
 - |— networkpolicy.yaml
 - |— rbac.yaml
 - |— servicemonitor.yaml
- |— tests/
 - |— test-connection.yaml
 - |— test-health.yaml

values-prod.yaml (Production Configuration)

global:

environment: production

imageTag: "v1.0.0"

imagePullPolicy: IfNotPresent

replicaCount: 3

image:

repository: gcr.io/sarthi-platform/sadp-api

tag: v1.0.0

service:

type: ClusterIP

port: 80

targetPort: 8000

ingress:

enabled: true

className: nginx

annotations:

cert-manager.io/cluster-issuer: letsencrypt-prod

nginx.ingress.kubernetes.io/rate-limit: "100"

nginx.ingress.kubernetes.io/ssl-redirect: "true"

hosts:

- host: sadp.sarthi.com

paths:

- path: /

pathType: Prefix

tls:

- secretName: sadp-tls

hosts:

- sadp.sarthi.com

autoscaling:

enabled: true

minReplicas: 3

maxReplicas: 20

targetCPUUtilizationPercentage: 70

targetMemoryUtilizationPercentage: 80

podDisruptionBudget:

enabled: true

minAvailable: 2

resources:

limits:

cpu: 2000m

memory: 4Gi

requests:

cpu: 500m

memory: 1Gi

nodeSelector:

node-type: compute-optimized

tolerations:

- key: "sadb-workload"

operator: "Equal"

value: "true"

effect: "NoSchedule"

affinity:

podAntiAffinity:

preferredDuringSchedulingIgnoredDuringExecution:

- weight: 100

podAffinityTerm:

labelSelector:

matchExpressions:

- key: app.kubernetes.io/name

operator: In

values:

- sadp-api

topologyKey: kubernetes.io/hostname

CI/CD Pipeline with GitLab CI

.gitlab-ci.yml: |

stages:

- test
- build
- security-scan
- deploy-staging
- integration-tests
- deploy-production
- post-deploy-tests

variables:

DOCKER_REGISTRY: gcr.io/sarathi-platform

DOCKER_DRIVER: overlay2

DOCKER_TLS_CERTDIR: "/certs"

before_script:

- echo \$GCP_SERVICE_ACCOUNT_KEY | base64 -d > gcp-key.json
- gcloud auth activate-service-account --key-file gcp-key.json
- gcloud config set project sarathi-platform

test:

stage: test

image: python:3.11

services:

- postgres:15
- redis:7

script:

- pip install poetry
- poetry install
- poetry run pytest tests/ --cov=src/ --cov-report=xml
- poetry run mypy src/
- poetry run black --check src/
- poetry run isort --check-only src/

coverage: '/TOTAL.+ ([0-9]{1,3}%)/'

artifacts:

reports:

coverage_report:

coverage_format: cobertura

path: coverage.xml

build:

stage: build

image: docker:20.10.16

services:

- docker:20.10.16-dind

script:

- docker build -t \$DOCKER_REGISTRY/sadp-api:\$CI_COMMIT_SHA .

- docker push \$DOCKER_REGISTRY/sadp-api:\$CI_COMMIT_SHA

only:

- main

- develop

security-scan:

stage: security-scan

image: aquasec/trivy:latest

script:

- trivy image --exit-code 1 --severity HIGH,CRITICAL \$DOCKER_REGISTRY/sadp-api:\$CI_COMMIT_SHA

allow_failure: false

deploy-staging:

stage: deploy-staging

image: google/cloud-sdk:alpine

script:

- gcloud container clusters get-credentials staging-cluster --region us-central1

- helm upgrade --install sadp-staging ./helm/sadp-platform

--namespace sadp-staging

--values ./helm/sadp-platform/values-staging.yaml

--set image.tag=\$CI_COMMIT_SHA

environment:

name: staging

url: https://sadp-staging.sarthi.com

only:

- develop

integration-tests:

stage: integration-tests

image: python:3.11

script:

- pip install pytest requests
- pytest tests/integration/ --base-url=https://sadb-staging.sarthi.com

dependencies:

- deploy-staging

deploy-production:

stage: deploy-production

image: google/cloud-sdk:alpine

script:

- gcloud container clusters get-credentials production-cluster --region us-central1
- helm upgrade --install sadp-production ./helm/sadb-platform
- namespace sadp-production
- values ./helm/sadb-platform/values-prod.yaml
- set image.tag=\$CI_COMMIT_SHA

environment:

name: production

url: https://sadb.sarthi.com

when: manual

only:

- main

post-deploy-tests:

stage: post-deploy-tests

image: python:3.11

script:





- pip install pytest requests
- pytest tests/smoke/ --base-url=https://sadb.sarthi.com

dependencies:

- deploy-production

Key Improvements Summary

1. Enterprise Architecture Patterns

-  **CQRS + Event Sourcing** for better scalability
-  **Service Mesh (Istio)** for advanced traffic management
-  **Circuit Breakers** for resilience
-  **Multi-level caching** for performance

2. Production-Grade Security

-  **Multi-factor authentication** for sensitive operations

- ☒ **Fine-grained authorization** with Open Policy Agent
- ☒ **Advanced encryption** with key rotation
- ☒ **Comprehensive audit logging** for HIPAA compliance

3. Advanced Observability

- ☒ **Distributed tracing** with Jaeger
- ☒ **Comprehensive metrics** with Prometheus
- ☒ **Intelligent alerting** with context awareness
- ☒ **Structured logging** with correlation IDs

4. Robust Error Handling

- ☒ **Intelligent error classification** with ML
- ☒ **Advanced retry patterns** with backoff strategies
- ☒ **Graceful degradation** mechanisms
- ☒ **Automatic recovery** procedures

5. Scalability & Performance

- ☒ **Connection pooling** for external services
- ☒ **Resource optimization** algorithms
- ☒ **Horizontal auto-scaling** based on metrics
- ☒ **Cost optimization** tracking and alerts

This enhanced architecture provides enterprise-grade capabilities that will scale with your Sarthi platform while maintaining the highest standards of security, compliance, and performance. It's designed to be implementation-ready for Claude CLI with clear specifications and comprehensive documentation.

Service Architecture

1. API Gateway Layer

FastAPI-based REST API with the following endpoints:

```
python
```

Agent Runtime APIs

POST /api/v1/agents/{agent_name}/execute
GET /api/v1/agents/{agent_name}/capabilities
POST /api/v1/workflows/execute
GET /api/v1/workflows/{workflow_id}/status

Evaluation APIs

POST /api/v1/evaluation/agents/{agent_name}/test
GET /api/v1/evaluation/agents/{agent_name}/results
POST /api/v1/evaluation/test-suites/generate
GET /api/v1/evaluation/benchmarks

Development APIs

POST /api/v1/development/prompts/deploy
POST /api/v1/development/prompts/ab-test
GET /api/v1/development/prompts/{prompt_id}/performance
POST /api/v1/development/ci-cd/trigger

Monitoring APIs

GET /api/v1/monitoring/agents/{agent_name}/metrics
POST /api/v1/monitoring/alerts/configure
GET /api/v1/monitoring/dashboard/{dashboard_id}
GET /api/v1/monitoring/performance/reports

2. Core Services

Agent Runtime Service

python

```

class AgentRuntimeService:
    """Core service for executing AI agents"""

    async def execute_agent(self, agent_name: str, input_data: dict,
                           context: dict = None) -> dict:
        """Execute specified agent with input data"""

    async def execute_workflow(self, workflow_definition: dict,
                              input_data: dict) -> dict:
        """Execute multi-agent workflow"""

    def get_agent_capabilities(self, agent_name: str) -> dict:
        """Get agent capabilities and metadata"""

    def register_agent(self, agent_config: dict) -> str:
        """Register new agent with the platform"""

```

Evaluation Service

```

python

class EvaluationService:
    """Service for agent evaluation and testing"""

    async def evaluate_agent(self, agent_name: str, test_cases: list) -> dict:
        """Run comprehensive agent evaluation"""

    async def generate_test_suite(self, agent_name: str,
                                 test_categories: list) -> dict:
        """Generate automated test suites"""

    def get_evaluation_results(self, evaluation_id: str) -> dict:
        """Get detailed evaluation results"""

    def compare_agent_versions(self, agent_name: str,
                              version_a: str, version_b: str) -> dict:
        """Compare performance between agent versions"""

```

Development Service

```

python

```

```

class DevelopmentService:
    """Service for agent development and deployment"""

    def deploy_prompt(self, prompt_config: dict) -> dict:
        """Deploy new prompt version"""

    async def setup_ab_test(self, test_config: dict) -> str:
        """Setup A/B testing for prompt optimization"""

    def trigger_ci_cd(self, agent_name: str, changes: dict) -> dict:
        """Trigger CI/CD pipeline for agent updates"""

    def get_deployment_status(self, deployment_id: str) -> dict:
        """Get deployment status and metrics"""

```

Monitoring Service

```

python

class MonitoringService:
    """Service for real-time monitoring and analytics"""

    def track_agent_execution(self, execution_data: dict) -> None:
        """Track agent execution metrics"""

    def get_performance_metrics(self, agent_name: str,
                                time_range: str) -> dict:
        """Get performance metrics for specified time range"""

    def configure_alerts(self, alert_config: dict) -> str:
        """Configure performance alerts"""

    def generate_analytics_report(self, report_config: dict) -> dict:
        """Generate comprehensive analytics reports"""

```

Implementation Stack

Backend Technology Stack

```

yaml

```


Core Framework

api_framework: "FastAPI 0.104+"

async_runtime: "Python 3.11+ with asyncio"

web_server: "Uvicorn with Gunicorn"

Database Layer

primary_db: "PostgreSQL 15+" # Agent configs, evaluations, metadata

cache_layer: "Redis 7+" # Session cache, real-time data

time_series: "InfluxDB 2.0" # Performance metrics, monitoring

document_store: "MongoDB 6+" # Test cases, reports, logs

Message Queue & Streaming

message_queue: "Apache Kafka" # Event streaming, workflow orchestration

task_queue: "Celery with Redis" # Background task processing

real_time: "WebSocket connections" # Live monitoring dashboards

Storage & CDN

file_storage: "Google Cloud Storage" # Prompt templates, test datasets

cdn: "Google Cloud CDN" # Static assets, cached responses

External APIs

ai_provider: "Anthropic Claude API"

healthcare_apis: "FHIR R4, HL7"

monitoring: "Google Cloud Monitoring"

logging: "Google Cloud Logging"

Security & Compliance

authentication: "OAuth 2.0 + JWT"

authorization: "RBAC with Casbin"

encryption: "AES-256 at rest, TLS 1.3 in transit"

compliance: "HIPAA, SOC 2 Type II"

Container Architecture

yaml

Docker Compose Structure

version: '3.8'

services:

API Gateway

sadp-gateway:

image: "nginx:alpine"

ports: ["80:80", "443:443"]

Core Services

sadp-api:

image: "sadp/api:latest"

replicas: 3

environment:

- DATABASE_URL=postgresql://...
- REDIS_URL=redis://...
- CLAUDE_API_KEY=\${CLAUDE_API_KEY}

sadp-worker:

image: "sadp/worker:latest"

replicas: 5

environment:

- CELERY_BROKER=redis://...

Data Layer

postgres:

image: "postgres:15-alpine"

environment:

- POSTGRES_DB=sadp
- POSTGRES_USER=sadp

volumes:

- postgres_data:/var/lib/postgresql/data

redis:

image: "redis:7-alpine"

volumes:

- redis_data:/data

mongodb:

image: "mongo:6"

volumes:

- mongo_data:/data/db

influxdb:

image: "influxdb:2.0"

volumes:

- influx_data:/var/lib/influxdb2

Message Queue

kafka:

image: "confluentinc/cp-kafka:latest"

environment:

- KAFKA_ZOOKEEPER_CONNECT=zookeeper:2181

zookeeper:

image: "confluentinc/cp-zookeeper:latest"

API Specifications

Agent Execution API

yaml

```
# POST /api/v1/agents/{agent_name}/execute
```

request:

agent_name: "clinical_agent"

input_data:

patient_id: "SARTHI-PT-001"

primary_diagnosis: "Type 2 Diabetes Mellitus"

patient_age: 55

comorbidities: ["Hypertension", "Obesity"]

context:

facility_id: "SARTHI-CLINIC-001"

provider_id: "DR-SMITH-001"

session_id: "sess_123456"

options:

timeout: 30000 # 30 seconds

priority: "high"

trace_enabled: true

response:

execution_id: "exec_789012"

agent_name: "clinical_agent"

status: "completed"

execution_time: 2847 # milliseconds

result:

treatment_plan: "Comprehensive treatment plan..."

confidence_score: 0.94

clinical_recommendations: [...]

metadata:

prompt_version: "v1.2.3"

claude_model: "claude-sonnet-4-20250514"

tokens_used: {"input": 1250, "output": 2890}

cost: 0.045 # USD

compliance:

hipaa_compliant: true

audit_trail_id: "audit_345678"

phi_processed: true

performance:

accuracy_score: 95.2

latency_percentile: "p95"

quality_metrics: {...}

Workflow Execution API

yaml

POST /api/v1/workflows/execute

request:

workflow_name: "patient_intake_complete"

input_data:

patient_id: "SARTHI-PT-002"

intake_form_image: "base64_encoded_image"

insurance_info: {...}

workflow_definition:

steps:

- agent: "document_processor"

input_mapping: {"document_image": "intake_form_image"}

output_key: "extracted_data"

- agent: "clinical_agent"

input_mapping: {"clinical_data": "extracted_data"}

output_key: "treatment_plan"

- agent: "billing_agent"

input_mapping: {"encounter_data": "extracted_data"}

output_key: "billing_claim"

parallel_steps: ["clinical_agent", "billing_agent"]

options:

timeout: 60000

priority: "normal"

response:

workflow_id: "wf_456789"

status: "completed"

total_execution_time: 8934 # milliseconds

steps_completed: 3

results:

document_processing:

execution_time: 4200

result: {...}

clinical_analysis:

execution_time: 3100

result: {...}

billing_verification:

execution_time: 1634

result: {...}

workflow_metrics:

total_cost: 0.127

accuracy_aggregate: 94.7

performance_rating: "excellent"

Evaluation API

yaml

POST /api/v1/evaluation/agents/{agent_name}/test

request:

agent_name: "clinical_agent"

test_suite_id: "clinical_comprehensive_v1"

evaluation_config:

test_categories: ["basic_functionality", "edge_cases", "compliance"]

performance_targets:

accuracy: 95

latency: 5000

compliance_score: 100

options:

parallel_execution: true

detailed_reporting: true

response:

evaluation_id: "eval_789123"

status: "in_progress"

estimated_completion: "2025-08-19T10:45:00Z"

test_cases_total: 47

test_cases_completed: 0

real_time_url: "wss://sadp.sarthi.com/evaluations/eval_789123/stream"

Monitoring API

yaml

```
# GET /api/v1/monitoring/agents/{agent_name}/metrics
```

request:

```
agent_name: "clinical_agent"
time_range: "24h"
metrics: ["accuracy", "latency", "error_rate", "cost"]
aggregation: "mean"
```

response:

```
agent_name: "clinical_agent"
time_range: "2025-08-18T10:00:00Z to 2025-08-19T10:00:00Z"
metrics:
  accuracy:
    current: 94.7
    trend: "stable"
    target: 95.0
    status: "warning"
  latency:
    current: 2847
    p95: 4200
    target: 5000
    status: "healthy"
  error_rate:
    current: 0.023
    trend: "improving"
    target: 0.05
    status: "healthy"
  cost:
    current: 0.045
    daily_total: 127.34
    trend: "increasing"
alerts:
  active: 1
  details: ["Accuracy below target threshold"]
recommendations:
  - "Consider prompt optimization to improve accuracy"
  - "Monitor cost trends for potential optimization"
```

Integration with Main Sarthi Application

SDK for Sarthi Application

```
python
```

Sarthi Application Integration SDK

class SarthiAISDK:

```
def __init__(self, api_base_url: str, api_key: str):
    self.api_base_url = api_base_url
    self.api_key = api_key
    self.session = httpx.AsyncClient()
```

async def execute_agent(self, agent_name: str, input_data: dict, **kwargs) -> dict:

```
    """Execute AI agent via SADP"""
    response = await self.session.post(
        f"{self.api_base_url}/api/v1/agents/{agent_name}/execute",
        json={"input_data": input_data, **kwargs},
        headers={"Authorization": f"Bearer {self.api_key}"})
    return response.json()
```

async def execute_workflow(self, workflow_name: str, input_data: dict) -> dict:

```
    """Execute multi-agent workflow"""
    response = await self.session.post(
        f"{self.api_base_url}/api/v1/workflows/execute",
        json={"workflow_name": workflow_name, "input_data": input_data},
        headers={"Authorization": f"Bearer {self.api_key}"})
    return response.json()
```

async def get_agent_performance(self, agent_name: str, time_range: str = "1h") -> dict:

```
    """Get real-time agent performance metrics"""
    response = await self.session.get(
        f"{self.api_base_url}/api/v1/monitoring/agents/{agent_name}/metrics",
        params={"time_range": time_range},
        headers={"Authorization": f"Bearer {self.api_key}"})
    return response.json()
```

Usage in Sarthi Application

class SarthiPatientService:

```
def __init__(self):
    self.ai_sdk = SarthiAISDK(
        api_base_url="https://sadp.sarthi.com",
        api_key=os.getenv("SADP_API_KEY"))
```

async def process_patient_intake(self, patient_data: dict) -> dict:


```
"""Process patient intake using SADP agents"""
```

```
# Execute patient intake workflow
```

```
workflow_result = await self.ai_sdk.execute_workflow(  
    workflow_name="patient_intake_complete",  
    input_data=patient_data  
)
```

```
# Extract results
```

```
document_analysis = workflow_result["results"]["document_processing"]  
clinical_summary = workflow_result["results"]["clinical_analysis"]  
billing_verification = workflow_result["results"]["billing_verification"]
```

```
return {  
    "patient_id": patient_data["patient_id"],  
    "intake_summary": clinical_summary,  
    "billing_ready": billing_verification,  
    "processing_time": workflow_result["total_execution_time"],  
    "ai_confidence": workflow_result["workflow_metrics"]["accuracy_aggregate"]  
}
```

```
async def get_ai_performance_dashboard(self) -> dict:
```

```
    """Get AI performance metrics for admin dashboard"""
```

```
agents = ["clinical_agent", "billing_agent", "document_processor"]  
performance_data = {}
```

```
for agent in agents:  
    metrics = await self.ai_sdk.get_agent_performance(agent, "24h")  
    performance_data[agent] = metrics
```

```
return performance_data
```

Deployment Architecture

Google Cloud Platform Deployment

```
yaml
```

GCP Services Architecture

gcp_services:

compute:

- service: "Google Kubernetes Engine (GKE)"
purpose: "Container orchestration for SADP services"
configuration:
 - node_pools: 3
 - auto_scaling: true
 - preemptible_nodes: true
- service: "Google Cloud Run"
purpose: "Serverless API endpoints for low-latency requests"
configuration:
 - concurrency: 100
 - timeout: 300s

storage:

- service: "Google Cloud SQL (PostgreSQL)"
purpose: "Primary database for agent configs and metadata"
configuration:
 - high_availability: true
 - automatic_backups: true
- service: "Google Cloud Storage"
purpose: "Prompt templates, test datasets, reports"
configuration:
 - storage_class: "STANDARD"
 - versioning: true

networking:

- service: "Google Cloud Load Balancer"
purpose: "API gateway and traffic distribution"
configuration:
 - ssl_termination: true
 - health_checks: true
- service: "Google Cloud CDN"
purpose: "Cache static assets and frequent responses"

monitoring:

- service: "Google Cloud Monitoring"
purpose: "Infrastructure and application monitoring"

- service: "Google Cloud Logging"
purpose: "Centralized logging and audit trails"

security:

- service: "Google Cloud IAM"
purpose: "Access control and service authentication"
- service: "Google Cloud KMS"
purpose: "Encryption key management"

Kubernetes Deployment Manifests

yaml

SADP API Deployment

apiVersion: apps/v1

kind: Deployment

metadata:

name: sadp-api

namespace: sadp

spec:

replicas: 3

selector:

matchLabels:

app: sadp-api

template:

metadata:

labels:

app: sadp-api

spec:

containers:

- name: sadp-api

image: gcr.io/sarhi-platform/sadp-api:v1.0.0

ports:

- containerPort: 8000

env:

- name: DATABASE_URL

valueFrom:

secretKeyRef:

name: sadp-secrets

key: database-url

- name: CLAUDE_API_KEY

valueFrom:

secretKeyRef:

name: sadp-secrets

key: claude-api-key

resources:

requests:

memory: "512Mi"

cpu: "250m"

limits:

memory: "1Gi"

cpu: "500m"

livenessProbe:

httpGet:

path: /health

port: 8000

```
    initialDelaySeconds: 30
    periodSeconds: 10
  readinessProbe:
    httpGet:
      path: /ready
      port: 8000
    initialDelaySeconds: 5
    periodSeconds: 5
```

SADP Service

```
apiVersion: v1
kind: Service
metadata:
  name: sadp-api-service
  namespace: sadp
spec:
  selector:
    app: sadp-api
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8000
  type: ClusterIP
```

SADP Ingress

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: sadp-ingress
  namespace: sadp
  annotations:
    kubernetes.io/ingress.class: "gce"
    cert-manager.io/cluster-issuer: "letsencrypt-prod"
spec:
  tls:
    - hosts:
        - sadp.sarthi.com
      secretName: sadp-tls
  rules:
    - host: sadp.sarthi.com
      http:
        paths:
```

```
- path: /api
pathType: Prefix
backend:
  service:
    name: sadp-api-service
    port:
      number: 80
```

Security & Compliance

Authentication & Authorization

```
python

# JWT-based authentication with RBAC
class SADPAuth:
    roles = {
        "sarathi_admin": [
            "agents:read", "agents:write", "agents:deploy",
            "evaluation:read", "evaluation:write",
            "monitoring:read", "development:read", "development:write"
        ],
        "sarathi_developer": [
            "agents:read", "evaluation:read", "evaluation:write",
            "development:read", "monitoring:read"
        ],
        "sarathi_application": [
            "agents:execute", "workflows:execute",
            "monitoring:read"
        ],
        "sarathi_viewer": [
            "agents:read", "evaluation:read", "monitoring:read"
        ]
    }

    def verify_permission(self, user_role: str, action: str) -> bool:
        return action in self.roles.get(user_role, [])
```

HIPAA Compliance Features

```
python
```

```

class HIPAACompliance:
    """HIPAA compliance enforcement for SADP"""

    def audit_log_interaction(self, request_data: dict) -> str:
        """Log all interactions for HIPAA audit trail"""
        audit_entry = {
            "timestamp": datetime.utcnow().isoformat(),
            "user_id": request_data.get("user_id"),
            "action": request_data.get("action"),
            "resource": request_data.get("resource"),
            "phi_accessed": self.detect_phi(request_data),
            "ip_address": request_data.get("ip_address"),
            "user_agent": request_data.get("user_agent")
        }
        return self.store_audit_log(audit_entry)

    def encrypt_sensitive_data(self, data: dict) -> dict:
        """Encrypt PHI and sensitive data"""
        # Implementation for AES-256 encryption
        pass

    def validate_data_access(self, user_id: str, patient_id: str) -> bool:
        """Validate user has permission to access patient data"""
        # Implementation for access control validation
        pass

```

Cost Optimization

Resource Management

```
python
```

```

class SADPResourceOptimizer:
    """Optimize resource usage and costs"""

    def auto_scale_workers(self, current_load: float) -> dict:
        """Auto-scale worker instances based on load"""
        if current_load > 0.8:
            return {"action": "scale_up", "instances": 2}
        elif current_load < 0.3:
            return {"action": "scale_down", "instances": 1}
        return {"action": "maintain", "instances": 0}

    def optimize_claude_api_usage(self, agent_metrics: dict) -> dict:
        """Optimize Claude API token usage"""
        recommendations = []

        if agent_metrics["avg_tokens"] > 3000:
            recommendations.append("Consider prompt length optimization")

        if agent_metrics["api_calls_per_hour"] > 1000:
            recommendations.append("Implement response caching")

        return {"recommendations": recommendations}

```

Benefits of Independent Service Architecture

1. Separation of Concerns

- **Main Sarthi App** focuses on healthcare workflows
- **SADP** focuses on AI agent management and optimization
- Clear boundaries and responsibilities

2. Independent Scaling

- Scale SADP based on AI workload demands
- Scale Sarthi app based on user traffic
- Optimize costs for each service independently

3. Technology Flexibility

- Use different tech stacks optimized for each purpose
- Upgrade AI capabilities without affecting main application

- Experiment with new technologies in isolation

4. Development Velocity

- Teams can develop and deploy independently
- Parallel development workflows
- Reduced merge conflicts and dependencies

5. Reliability & Fault Isolation

- SADP failures don't bring down main application
- Independent monitoring and alerting
- Easier debugging and troubleshooting

6. Security & Compliance

- Dedicated security controls for AI operations
- Isolated compliance auditing
- Reduced attack surface for main application

7. Cost Optimization

- Pay only for AI resources when needed
- Independent cost tracking and optimization
- Better resource utilization

This independent service architecture provides maximum flexibility, scalability, and maintainability while enabling your main Sarthi application to leverage advanced AI capabilities through a clean, well-defined API interface.